

## [Hbase – 2.4]

1.	pre-requisite .....	3
2.	Hbase Installation – 60 Minutes .....	4
3.	Using the HBase Shell & Data Access – 60 minutes .....	7
4.	Understanding Hbase Architecture – 30 Minutes .....	19
5.	Hadoop Single Node Cluster – 60 Minutes .....	30
6.	Using HDFS – 30 Minutes .....	41
7.	Pseudo Distributed Hbase Cluster – 60 Minutes .....	48
8.	Using the Developer JAVA API - 60 Minutes .....	53
9.	HBase Filters – 30 Minutes .....	77
10.	Built-in TSV Bulk Loader – 60 Minutes .....	85
11.	Hbase Spark Integration – Spark-Hbase connector – 90 Minutes .....	98
12.	Install MySQL on a CentOS 7 Server..... Start MySQL and Check its Status .....	108
13.	Errors: .....	109
14.	Annexure: .....	113
15.	Annexure (Standalone Mode) .....	114

## [Hbase – 2.4]

Last Updated: 16<sup>th</sup> May 2022.

Hbase with Phoenix:

Hadoop Version : 3.1.2 , Hbase Version : 2.4.11 and Phoenix : 5.1.2 (phoenix-hbase-2.4.0-5.1.2-bin.tar.gz)

Hbase cluster only:

Hbase version : 3.1.2 and Hadoop : 3.1.2

<https://community.cloudera.com/t5/Community-Articles/Compression-in-HBase/ta-p/247244>

## 1. pre-requisite

You can install Hbase either in VM or in Docker container.

1) VM

Start your VM. Use Centos or Redhat Linux version above 6.0

2) On Docker:

Only Hbase – Cluster

```
# docker run --name hbaseo --hostname hmaster.master.com --privileged -p 2181:2181 -p 16000:16000 -p 16010:16010 -p 16002:16002 -p 16012:16012 -p 16020:16020 -p 16030:16030 -p 8080:8080 -p 8085:8085 -p 9090:9090 -p 9095:9095 -i -t centos:7 /usr/sbin/init
```

Hbase & YARN Cluster

```
#docker run --name hbaseo --hostname hmaster.master.com --privileged -p 2181:2181 -p 16000:16000 -p 16010:16010 -p 16002:16002 -p 16012:16012 -p 16020:16020 -p 16030:16030 -p 8080:8080 -p 8085:8085 -p 8088:8088 -p 9870:9870 -p 9864:9864 -i -t centos:7 /usr/sbin/init
```

## 2. Hbase Installation – 60 Minutes

HBase requires that a JDK be installed. Use jdk - [jdk-8u121-linux-x64.tar.gz](#)

Extract it and specify the JAVA\_HOME in `~/.bashrc` as mention below.

```
export PATH=$PATH:/opt/jdk/bin:/opt/hbase/bin  
export JAVA_HOME=/opt/jdk
```

Extract the downloaded file, and change to the newly-created directory.

```
$ tar xzvf hbase-3.0.0-SNAPSHOT-bin.tar.gz -C /opt  
#mv hbase-3.0.0-SNAPSHOT-bin hbase  
$ cd hbase/
```

You must set the `JAVA_HOME` environment variable before starting HBase. To make this easier, HBase lets you set it within the `conf/hbase-env.sh` file.

```
# Set environment variables here.  
# The java implementation to use.  
export JAVA_HOME=/opt/jdk
```

The `bin/start-hbase.sh` script is provided as a convenient way to start HBase.

Issue the command, and if all goes well, a message is logged to standard output showing that HBase started successfully.

## [Hbase – 2.4]

```
[root@4f5607f478bd hbase]# bin/start-hbase.sh
running master, logging to /opt/hbase/bin/../logs/hbase--master-4f5607f478bd.out
[root@4f5607f478bd hbase]# tail -f logs/hbase--master-4f5607f478bd.out
WARNING: An illegal reflective access operation has occurred
WARNING: Illegal reflective access by org.apache.hadoop.hbase.util.UnsafeAvailChecker (file:/opt/hbase/lib/hbase-common-2.3.2.jar) to method
java.nio.Bits.unaligned()
WARNING: Please consider reporting this to the maintainers of org.apache.hadoop.hbase.util.UnsafeAvailChecker
WARNING: Use --illegal-access=warn to enable warnings of further illegal reflective access operations
WARNING: All illegal access operations will be denied in a future release
```

Check the log in case of any issue.

```
0x100000ae1d2000
2020-11-25 08:00:16,204 INFO [ReadOnlyZKClient-127.0.0.1:2181@0x5943051a] zookeeper.ZooKeeper: Session: 0x100000ae1d20002 closed
2020-11-25 08:00:16,205 INFO [ReadOnlyZKClient-127.0.0.1:2181@0x5943051a-EventThread] zookeeper.ClientCnxn: EventThread shut down for session: 0x100000ae1d20002
[root@4f5607f478bd logs]# ls
SecurityAuth.audit  hbase--master-4f5607f478bd.log  hbase--master-4f5607f478bd.out  hbase--master-4f5607f478bd.out.1
[root@4f5607f478bd logs]# pwd
/opt/hbase/logs
[root@4f5607f478bd logs]#
```

You can use the jps command to verify that you have one running process called HMaster.

```
/opt/hbase/logs
[root@4f5607f478bd logs]# jps
626 Jps
199 HMaster
[root@4f5607f478bd logs]#
```

In standalone mode HBase runs all daemons within this single JVM, i.e. the HMaster, a single HRegionServer, and the ZooKeeper daemon.

Go to <http://localhost:16010> to view the HBase Web UI

OSTech

[Hbase – 2.4]

The screenshot shows the Apache HBase master interface. At the top, there is a navigation bar with links: Home, Table Details, Procedures & Locks, HBCK Report, Process Metrics, Local Logs, Log Level, Debug Dump, Metrics Dump, Profiler, and HBase Configuration. Below the navigation bar, it says "MASTER 4f5607f478bd". The main content area is titled "Region Servers". There is a table with the following data:

Base Stats	Memory	Requests	Storefiles	Compactions	Replications			
ServerName				Start time	Last contact	Version	Requests Per Second	Num. Regions
4f5607f478bd,16020,1606291148731				Wed Nov 25 07:59:08 UTC 2020	2 s	2.3.2	0	2
Total:1							0	2

----- Lab Ends Here -----

OSTech

### **3. Using the HBase Shell & Data Access – 60 minutes**

*Procedure: Use HBase For the First Time*

1. Connect to HBase.

Connect to your running instance of HBase using the `hbase shell` command, located in the `bin/` directory of your HBase install. In this example, some usage and version information that is printed when you start HBase Shell has been omitted. The HBase Shell prompt ends with a `>` character.

```
$ ./bin/hbase shell
```

2. Display HBase Shell Help Text.

Type `help` and press Enter, to display some basic usage information for HBase Shell, as well as several example commands. Notice that table names, rows, columns all must be enclosed in quote characters.

## [Hbase – 2.4]

```
ses where applicable
HBase Shell
Use "help" to get list of supported commands.
Use "exit" to quit this interactive shell.
For Reference, please visit: http://hbase.apache.org/2.0/book.html#shell
Version 2.3.2, r8dc7de1c75ab615b8992b928150e8b476cb053e3, Sun Sep 20 13:28:26 UTC 2020
Took 0.0022 seconds
hbase(main):001:0> help
HBase Shell, version 2.3.2, r8dc7de1c75ab615b8992b928150e8b476cb053e3, Sun Sep 20 13:28:26 UTC 2020
Type 'help "COMMAND"', (e.g. 'help "get"' -- the quotes are necessary) for help on a specific command.
Commands are grouped. Type 'help "COMMAND_GROUP"', (e.g. 'help "general"') for help on a command group.

COMMAND GROUPS:
Group name: general
Commands: processlist, status, table_help, version, whoami
Group name: ddl
Commands: alter, alter_async, alter_status, clone_table_schema, create, describe, disable, disable_all, drop, drop_all, enable, enable_all, exists, get_table, is_disabled, is_enabled, list, list_regions, locate_region, show_filters
```

### 3. Create a table.

Use the `create` command to create a new table. You must specify the table name and the ColumnFamily name.

```
hbase(main):001:0> create 'test', 'cf'
```

### 4. List Information About your Table

Use the `list` command to confirm your table exists

```
hbase(main):002:0> list 'test'
```

## [Hbase – 2.4]

Now use the `describe` command to see details, including configuration defaults

```
hbase(main):003:0> describe 'test'
```

```
hbase(main):006:0> create 'test', 'cf'
Created table test
Took 1.2857 seconds
=> Hbase::Table - test
hbase(main):007:0> list 'test'
TABLE
test
1 row(s)
Took 0.0703 seconds
=> ["test"]
hbase(main):008:0>
hbase(main):009:0* describe 'test'
Table test is ENABLED
test
COLUMN FAMILIES DESCRIPTION
{NAME => 'cf', VERSIONS => '1', EVICT_BLOCKS_ON_CLOSE => 'false', KEEP_DELETED_CELLS => 'FALSE', CACHE_DATA_ON_WRITE => 'false', DATA_BLOCK_ENCODING => 'NONE', TTL => 'FOREVER', MIN_VERSIONS => '0', REPLICATION_SCOPE => '0', BLOOMFILTER => 'ROW', CACHE_INDEX_ON_WRITE => 'false', IN_MEMORY => 'false', CACHE_BLOOMS_ON_WRITE => 'false', PREFETCH_BLOCKS_ON_OPEN => 'false', COMPRESSION => 'NONE', BLOCKCACHE => 'true', BLOCK_SIZE => '65536'}
```

1 row(s)

QUOTAS

0 row(s)

Took 0.4240 seconds

## [Hbase – 2.4]

### 5. Put data into your table.

To put data into your table, use the `put` command.

```
hbase(main):003:0> put 'test', 'row1', 'cf:a', 'value1'  
hbase(main):004:0> put 'test', 'row2', 'cf:b', 'value2'  
hbase(main):005:0> put 'test', 'row3', 'cf:c', 'value3'
```

Here, we insert three values, one at a time. The first insert is at `row1`, column `cf:a`, with a value of `value1`. Columns in HBase are comprised of a column family prefix, `cf` in this example, followed by a colon and then a column qualifier suffix, `a` in this case.

### 6. Scan the table for all data at once.

One of the ways to get data from HBase is to scan. Use the `scan` command to scan the table for data. You can limit your scan, but for now, all data is fetched.

```
hbase(main):006:0> scan 'test'
```

## [Hbase – 2.4]

```
hbase(main):010:0>
hbase(main):011:0* put 'test', 'row1', 'cf:a', 'value1'
Took 0.0958 seconds
hbase(main):012:0> put 'test', 'row2', 'cf:b', 'value2'
Took 0.0077 seconds
hbase(main):013:0> put 'test', 'row3', 'cf:c', 'value3'
Took 0.0079 seconds
hbase(main):014:0> scan 'test'
ROW                                COLUMN+CELL
  row1                            column=cf:a, timestamp=2020-11-25T08:18:25.387, value=value1
  row2                            column=cf:b, timestamp=2020-11-25T08:18:31.324, value=value2
  row3                            column=cf:c, timestamp=2020-11-25T08:18:37.532, value=value3
3 row(s)
Took 0.0746 seconds
```

### 7. Get a single row of data.

To get a single row of data at a time, use the `get` command.

```
hbase(main):007:0> get 'test', 'row1'
```

### 8. Disable a table.

If you want to delete a table or change its settings, as well as in some other situations, you need to disable the table first, using the `disable` command. You can re-enable it using the `enable` command.

```
hbase(main):008:0> disable 'test'
```

Try fetching information from the table. It should throw an exception.

```
hbase(main):007:0> get 'test', 'row1'
```

## [Hbase – 2.4]

```
hbase(main):015:0> get 'test', 'row1'
COLUMN          CELL
cf:a           timestamp=2020-11-25T08:18:25.387, value=value1
1 row(s)
Took 0.0681 seconds
hbase(main):016:0> disable 'test'
Took 0.4304 seconds
hbase(main):017:0> get 'test', 'row1'
COLUMN          CELL

ERROR: Table test is disabled!
For usage try 'help "get"'
Took 0.4916 seconds

hbase(main):009:0> enable 'test'
```

Disable the table again if you tested the `enable` command above:

```
hbase(main):010:0> disable 'test'
```

### 9. Drop the table.

To drop (delete) a table, use the `drop` command.

```
hbase(main):011:0> drop 'test'
```

## [Hbase – 2.4]

```
hbase(main):018:0> enable 'test'
Took 0.6932 seconds
hbase(main):019:0> get 'test', 'row1'
COLUMN                                CELL
cf:a                                  timestamp=2020-11-25T08:18:25.387, value=value1
1 row(s)
Took 0.0701 seconds
hbase(main):020:0> disable 'test'
Took 0.3778 seconds
hbase(main):021:0> drop 'test'
Took 0.1648 seconds
```

10. Exit the HBase Shell.

To exit the HBase Shell and disconnect from your cluster, use the `quit` command. HBase is still running in the background.

## [Hbase – 2.4]

Start the Shell :

```
#bin/hbase shell  
create_namespace 'entertainment'
```

```
create 'entertainment:movie', {NAME => 'desc', VERSIONS => 2}, {NAME => 'media'}
```

```
list
```

```
hbase(main):001:0> create_namespace 'entertainment'  
Took 0.7972 seconds  
hbase(main):002:0> create 'entertainment:movie', {NAME => 'desc', VERSIONS => 2}, {NAME => 'media'}  
Created table entertainment:movie  
Took 0.6836 seconds  
=> Hbase::Table - entertainment:movie  
hbase(main):003:0> list  
TABLE  
entertainment:movie  
1 row(s)  
Took 0.0406 seconds  
=> ["entertainment:movie"]  
hbase(main):004:0> █
```

```
alter 'entertainment:movie', NAME => 'desc', VERSIONS => 5
```

OSTech

## [Hbase – 2.4]

```
put 'entertainment:movie', 'row1', 'desc:title', 'Home Alone'  
put 'entertainment:movie', 'row1', 'desc:title', 'Home Alone 2', 1274032629663  
put 'entertainment:movie', 'row2', 'desc:title', 'Braveheart'
```

```
hbase(main):005:0> alter 'entertainment:movie', NAME => 'desc', VERSIONS => 5  
Updating all regions with the new schema...  
1/1 regions updated.  
Done.  
Took 1.8517 seconds  
hbase(main):006:0> put 'entertainment:movie', 'row1', 'desc:title', 'Home Alone'  
Took 0.1570 seconds  
hbase(main):007:0> put 'entertainment:movie', 'row1', 'desc:title', 'Home Alone 2', 1274032629663  
Took 0.0105 seconds  
hbase(main):008:0> put 'entertainment:movie', 'row2', 'desc:title', 'Braveheart'  
Took 0.0079 seconds  
hbase(main):009:0> █
```

Different ways of retrieving row information in the Hbase:

```
get 'entertainment:movie', 'row1'
```

OSTech

## [Hbase – 2.4]

```
get 'entertainment:movie', 'row1', {COLUMN => 'desc:title'}  
get 'entertainment:movie', 'row1', {COLUMN => 'desc:title', VERSIONS => 2}  
get 'entertainment:movie', 'row1', {COLUMN => ['desc']}
```

[Hbase – 2.4]

```
hbase(main):010:0> get 'entertainment:movie', 'row1'
COLUMN                           CELL
  desc:title                     timestamp=2020-11-25T08:41:39.680, value=Home Alone
1 row(s)
Took 0.0560 seconds
hbase(main):011:0> get 'entertainment:movie', 'row1', {COLUMN => 'desc:title'}
COLUMN                           CELL
  desc:title                     timestamp=2020-11-25T08:41:39.680, value=Home Alone
1 row(s)
Took 0.0174 seconds
hbase(main):012:0> get 'entertainment:movie', 'row1', {COLUMN => 'desc:title', VERSIONS => 2}
COLUMN                           CELL
  desc:title                     timestamp=2020-11-25T08:41:39.680, value=Home Alone
  desc:title                     timestamp=2010-05-16T17:57:09.663, value=Home Alone 2
1 row(s)
Took 0.0340 seconds
hbase(main):013:0> get 'entertainment:movie', 'row1', {COLUMN => ['desc']}
COLUMN                           CELL
  desc:title                     timestamp=2020-11-25T08:41:39.680, value=Home Alone
1 row(s)
Took 0.0280 seconds
hbase(main):014:0> get 'entertainment:movie', 'row1', {COLUMN => 'desc:title', VERSIONS => 3}
COLUMN                           CELL
  desc:title                     timestamp=2020-11-25T08:41:39.680, value=Home Alone
  desc:title                     timestamp=2010-05-16T17:57:09.663, value=Home Alone 2
```

## [Hbase – 2.4]

Different ways of fetching rows :

```
scan 'entertainment:movie', {LIMIT => 10}
scan 'entertainment:movie', {STARTROW => 'row1',STOPROW => 'row5'}
scan 'entertainment:movie', {COLUMNS =>['desc:title', 'media:type']}
deleteall 'entertainment:movie', 'row1'
```

```
hbase(main):019:0> scan 'entertainment:movie', {LIMIT => 10}
ROW                                COLUMN+CELL
  row1                            column=desc:title, timestamp=2020-11-25T08:41:39.680, value=Home Alone
  row2                            column=desc:title, timestamp=2020-11-25T08:42:25.244, value=Braveheart
2 row(s)
Took 0.0478 seconds
hbase(main):020:0> scan 'entertainment:movie', {STARTROW => 'row1',STOPROW => 'row5'}
ROW                                COLUMN+CELL
  row1                            column=desc:title, timestamp=2020-11-25T08:41:39.680, value=Home Alone
  row2                            column=desc:title, timestamp=2020-11-25T08:42:25.244, value=Braveheart
2 row(s)
Took 0.0192 seconds
hbase(main):021:0> scan 'entertainment:movie', {COLUMNS =>['desc:title', 'media:type']}
ROW                                COLUMN+CELL
  row1                            column=desc:title, timestamp=2020-11-25T08:41:39.680, value=Home Alone
  row2                            column=desc:title, timestamp=2020-11-25T08:42:25.244, value=Braveheart
2 row(s)
Took 0.0220 seconds
hbase(main):022:0> deleteall 'entertainment:movie', 'row1'
Took 0.0086 seconds
hbase(main):023:0> scan 'entertainment:movie', {COLUMNS =>['desc:title', 'media:type']}
ROW                                COLUMN+CELL
  row2                            column=desc:title, timestamp=2020-11-25T08:42:25.244, value=Braveheart
1 row(s)
Took 0.0202 seconds
hbase(main):024:0>
```

----- Lab Ends Here -----

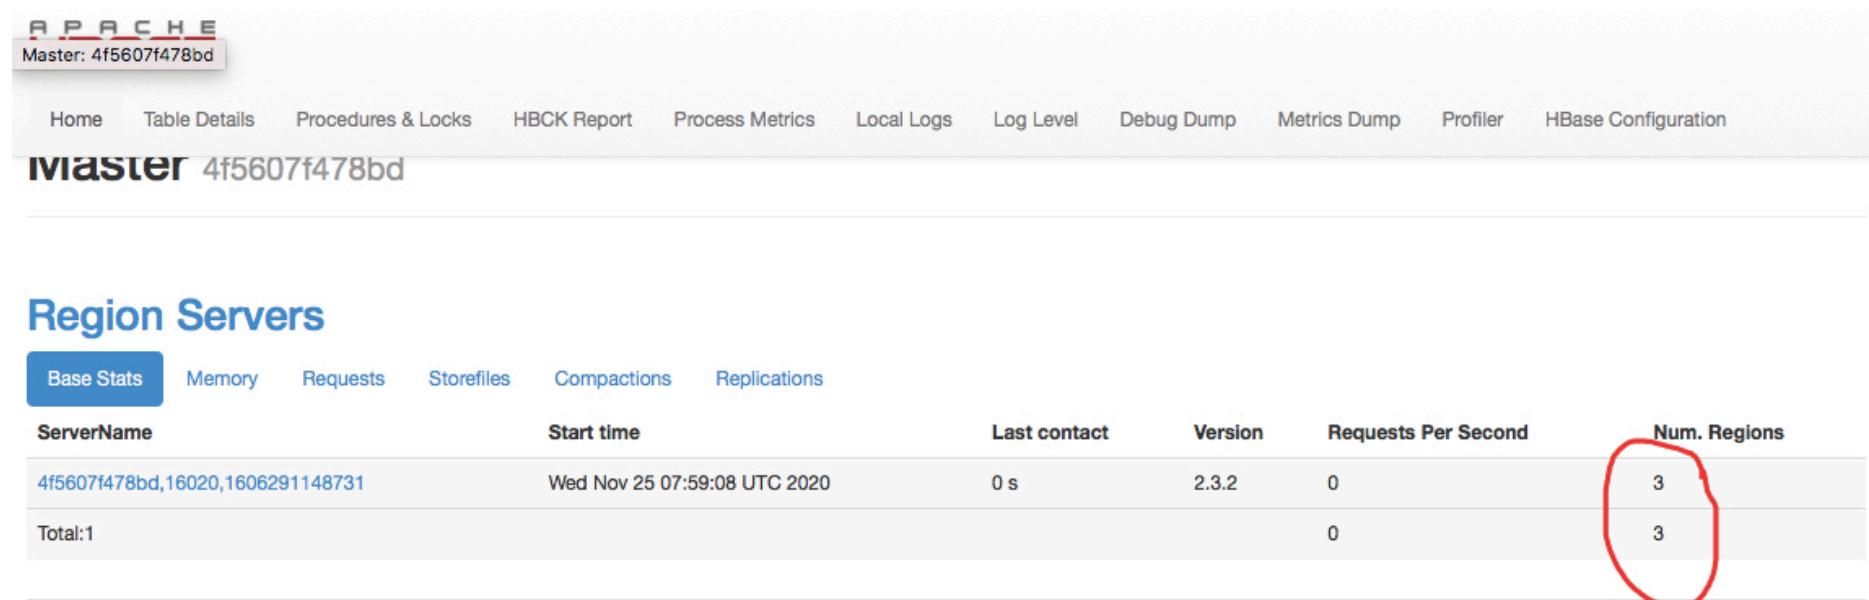
[Hbase – 2.4]

## 4. Understanding Hbase Architecture – 30 Minutes

Determine the no of region in a region Server.

Using Web UI.:

Determine the No of regions in our Existing Hbase Node.



Region Servers							
Base Stats	Memory	Requests	Storefiles	Compactions	Replications		
ServerName	Start time			Last contact	Version	Requests Per Second	Num. Regions
4f5607f478bd,16020,1606291148731	Wed Nov 25 07:59:08 UTC 2020			0 s	2.3.2	0	3
Total:1						0	3

In the above example, its having 3 regions in the node, 4f\*.

OSTech

## [Hbase – 2.4]

How many tables are there? Scroll down till you find a section as shown below in the Admin web UI.

Click on the User Tables Tab : There will be a table, movie which we have created earlier.

### Tables

			Regions								Description
Namespace	Name	State	OPEN	OPENING	CLOSED	CLOSING	OFFLINE	FAILED	SPLIT	Other	
			OPEN	OPENING	CLOSED	CLOSING	OFFLINE	FAILED	SPLIT	Other	
entertainment	movie	ENABLED	1	0	0	0	0	0	0	0	'entertainment:movie', {NAME => 'desc', VERSIONS => '5'}, {NAME => 'media'}

Click on the System Tables.

## [Hbase – 2.4]

The screenshot shows the HBase master interface. At the top, there's a header bar with the title "Tables" and a master identifier "Master: 4f5607f478bd". Below the header, there are three tabs: "User Tables", "System Tables" (which is currently selected and highlighted in blue), and "Snapshots". The main content area displays a table with two rows. The first row has columns "Table Name" and "Description". The second row contains the entries "hbase:meta" and "The hbase:meta table holds references to all User Table regions.", and "hbase:namespace" and "The hbase:namespace table holds information about namespaces.".

Table Name	Description
hbase:meta	The hbase:meta table holds references to all User Table regions.
hbase:namespace	The hbase:namespace table holds information about namespaces.

There are two system tables; meta and namespace.

Click on the Local Logs to determine the various Logs.



[Hbase – 2.4]

## Directory: /logs/

Name ▲	Last Modified	Size
hbase--master-4f5607f478bd.log	Nov 25, 2020, 9:14:56 AM	223,949 bytes
hbase--master-4f5607f478bd.out	Nov 25, 2020, 7:59:04 AM	514 bytes
hbase--master-4f5607f478bd.out.1	Oct 22, 2020, 3:51:43 PM	514 bytes
SecurityAuth.audit	Nov 25, 2020, 9:14:14 AM	7,335 bytes

Click on the hbase-\* .log to view its content.

## [Hbase – 2.4]

```
2020-11-25 09:09:05,011 INFO [LruBlockCacheStatsExecutor] hfile.LruBlockCache: totalSize=156.93 KB, freeSize=199.05 MB, max=199.20 MB, blockCount=5, accesses=53, hits=42, hitRatio=70.35%, cachingAccesses=47, cachingHits=38, cachingHitsRatio=80.85%, evictions=419, evicted=0, evictedPerRun=0.0
2| 127.0.0.1:16010/logs/hbase-->[master-4f5607f478bd:16000.Chore.1] balancer.BaseLoadBalancer: Start Generate Balance plan for cluster.
2| master-4f5607f478bd.log |FO [master/4f5607f478bd:16000.Chore.1] master.HMaster: Normalizer ran successfully in PT0.011S. Submitted 0 plans, affecting 0 tables.
2020-11-25 09:12:27,429 INFO [ReadOnlyZKClient-127.0.0.1:2181@0x066a7829] zookeeper.ZooKeeper: Initiating client connection, connectString=127.0.0.1:2181
sessionTimeout=10000 watcher=org.apache.hadoop.hbase.zookeeper.ReadOnlyZKClient$$Lambda$216/0x000000800f4d440@7c8c7741
2020-11-25 09:12:27,431 INFO [ReadOnlyZKClient-127.0.0.1:2181@0x066a7829] zookeeper.ClientCnxnSocket: jute.maxbuffer value is 4194304 Bytes
2020-11-25 09:12:27,431 INFO [ReadOnlyZKClient-127.0.0.1:2181@0x066a7829] zookeeper.ClientCnxn: zookeeper.request.timeout value is 0. feature.enabled=
2020-11-25 09:12:27,433 INFO [ReadOnlyZKClient-127.0.0.1:2181@0x066a7829-SendThread(127.0.0.1:2181)] zookeeper.ClientCnxn: Opening socket connection to server
localhost/127.0.0.1:2181. Will not attempt to authenticate using SASL (unknown error)
2020-11-25 09:12:27,435 INFO [ReadOnlyZKClient-127.0.0.1:2181@0x066a7829-SendThread(127.0.0.1:2181)] zookeeper.ClientCnxn: Socket connection established, initiating
session, client: /127.0.0.1:35672, server: localhost/127.0.0.1:2181
2020-11-25 09:12:27,438 INFO [ReadOnlyZKClient-127.0.0.1:2181@0x066a7829-SendThread(127.0.0.1:2181)] zookeeper.ClientCnxn: Session establishment complete on server
localhost/127.0.0.1:2181, sessionid = 0x100000ae1d2000d, negotiated timeout = 10000
2020-11-25 09:13:27,512 INFO [ReadOnlyZKClient-127.0.0.1:2181@0x066a7829] zookeeper.ZooKeeper: Session: 0x100000ae1d2000d closed
2020-11-25 09:13:27,513 INFO [ReadOnlyZKClient-127.0.0.1:2181@0x066a7829-EventThread] zookeeper.ClientCnxn: EventThread shut down for session: 0x100000ae1d2000d
2020-11-25 09:14:04,742 INFO [LruBlockCacheStatsExecutor] hfile.LruBlockCache: totalSize=156.93 KB, freeSize=199.05 MB, max=199.20 MB, blockCount=5, accesses=56, hits=45,
hitRatio=80.36%, cachingAccesses=50, cachingHits=41, cachingHitsRatio=82.00%, evictions=449, evicted=0, evictedPerRun=0.0
2020-11-25 09:14:14,048 INFO [master/4f5607f478bd:16000.Chore.1] master.HMaster: Normalizer ran successfully in PT0.012S. Submitted 0 plans, affecting 0 tables.
2020-11-25 09:14:14,052 INFO [master-4f5607f478bd:16000.Chore.1] balancer.BaseLoadBalancer: Start Generate Balance plan for cluster.
2020-11-25 09:14:56,034 INFO [master:store-WAL-Roller] wal.AbstractFSWAL: Rolled WAL
/opt/hbase/tmp/hbase/MasterData/WALs/4f5607f478bd,16000,1606291144491/4f5607f478bd@2C16000@2C1606291144491.1606294786757 with entries=0, filesize=83 B; new WAL
/opt/hbase/tmp/hbase/MasterData/WALs/4f5607f478bd,16000,1606291144491/4f5607f478bd@2C16000@2C1606291144491.1606295696011
2020-11-25 09:14:56,039 INFO [master:store-WAL-Roller] wal.AbstractFSWAL: Archiving
file:/opt/hbase/tmp/hbase/MasterData/WALs/4f5607f478bd,16000,1606291144491/4f5607f478bd@2C16000@2C1606291144491.1606294786757 to
file:/opt/hbase/tmp/hbase/MasterData/oldWALs/4f5607f478bd@2C16000@2C1606291144491.1606294786757
2020-11-25 09:14:56,042 INFO [master:store-WAL-Roller] region.MasterRegionUtils: Moved
file:/opt/hbase/tmp/hbase/MasterData/oldWALs/4f5607f478bd@2C16000@2C1606291144491.1606294786757 to
file:/opt/hbase/tmp/hbase/oldWALs/4f5607f478bd@2C16000@2C1606291144491.1606294786757$masterlocalwal$
```

You can look for errors in this file, if you are having issue.

As shown above, the current log level is INFO. You can determine the current log level using the Log Level tab in the Admin Web UI section.

Log Name :**hbase--master-4f5607f478bd.log** Click on Get Log Level.

## [Hbase – 2.4]

The screenshot shows the Apache HBase Log Level interface. At the top, there are several tabs: Email: Inbox (137), Trello, Apache HBase™ Reference..., Log Level (which is active and highlighted with a red circle), Edit Pad - Online Text Editor..., and Ruel - Hard Sometim... Below the tabs, the Apache HBase logo is displayed. The main navigation menu includes Home, Table Details, Procedures & Locks, HBCK Report, Process Metrics, Local Logs, Log Level (again), Debug Dump, Metrics Dump, Profiler, and HBase Configuration. The title of the page is "Get/Set Log Level".

**Actions:**

- Get Log Level** (button) ✓: Get the current log level for the specified log name.
- Set Log Level** (button): Set the specified log level for the specified log name. This section contains two input fields: "Log Name (required)" containing "hbase--master-4f5607f478bd.log" and "Log Level (required)" containing "INFO".

### Results:

Submitted Log Name: hbase--master-4f5607f478bd.log  
Log Class: org.slf4j.impl.Log4JLoggerAdapter  
Effective level: INFO

Its showing the current log level to INFO.  
Let us change it to WARNING.

hbase--master-4f5607f478bd.log  
DEBUG

Click on the Set LOG LEVEL.

OSTech

## [Hbase – 2.4]

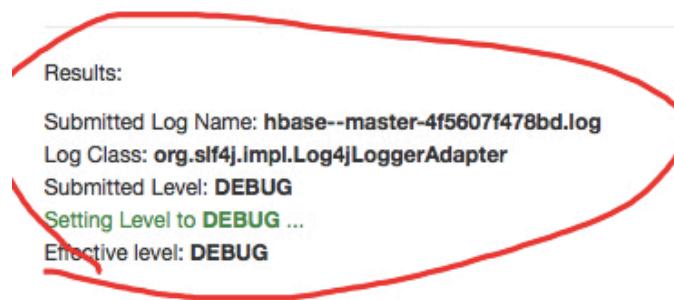
Log Level

Get Log Level      Log Name (required)      Get the current log level for the specified log name.

Set Log Level      hbase--master-4f5607f478bd.log      Set the specified log level for the specified log name.

Results:

Submitted Log Name: hbase--master-4f5607f478bd.log  
Log Class: org.slf4j.impl.Log4jLoggerAdapter  
Submitted Level: DEBUG  
Setting Level to DEBUG ...  
Effective level: DEBUG



It may take a few minutes.

Later verify the log level.

## [Hbase – 2.4]

The screenshot shows the Apache HBase Log Level interface. At the top, there is a navigation bar with links: Home, Table Details, Procedures & Locks, HBCK Report, Process Metrics, Local Logs, Log Level (which is highlighted), Debug Dump, Metrics Dump, Profiler, and HBase Configuration. Below the navigation bar, the title "Get/Set Log Level" is displayed. Under the "Actions:" section, there are two main buttons: "Get Log Level" (highlighted with a red underline) and "Set Log Level". To the right of the "Get Log Level" button is a text input field containing the value "hbase--master-4f5607f478bd.log", which is also highlighted with a red underline. A tooltip for this field states: "Get the current log level for the specified log name." Below these actions, under the "Results:" section, the following information is displayed:  
Submitted Log Name: hbase--master-4f5607f478bd.log  
Log Class: org.slf4j.impl.Log4JLoggerAdapter  
Effective level: DEBUG

It should be DEBUG as shown above.

You can find a region for a particular table using the following option:

Click on table Details → Select the User tables e.x [entertainment:movie](#)

## [Hbase – 2.4]

Table Regions											
Base Stats	Compactions	Region Server	ReadRequests (0)	WriteRequests (0)	StorefileSize (0 MB)	Num.Storefiles (2)	MemSize (0 MB)	Locality	Start Key	End Key	Region State
Name(1) entertainment:movie,,1606293502017.9a987ce29928e84840835b50 f0e90765.	4f5607f478bd:16030	0	0	0 MB	2	0 MB	1.0				OPEN

## Regions by Region Server

Region Server	Region Count
4f5607f478bd:16030	1

As shown above, there is a single region for this table which is open for request.

Understand the system tables:

Execute the following commands in the hbase shell.

```
describe 'hbase:meta'  
scan 'hbase:meta'
```

OSTech

[Hbase – 2.4]

records of movie and namespace tables are stored in the meta table as shown above.

[Hbase – 2.4]

----- Lab Ends Here -----

[Hbase – 2.4]

## 5. Hadoop Single Node Cluster – 60 Minutes

Hadoop : 3.1.2

JDK : jdk-8u221

<https://www.apache.org/dyn/closer.cgi/hadoop/common/hadoop-3.1.4/hadoop-3.1.4.tar.gz>

Extract in a folder and rename the folder as Hadoop as show below.

```
Command: None [opt]# pwd  
/opt  
[root@hmaster opt]# ls  
hadoop hadoop-3.1.4.tar hbase hbase-tools jdk  
[root@hmaster opt]#
```

### Prepare to Start the Hadoop Cluster

Unpack the downloaded Hadoop distribution. In the distribution, edit the file /opt/hadoop/etc/hadoop/hadoop-env.sh to define some parameters as follows or in the ~/.bashrc file:

```
export HDFS_NAMENODE_USER="root"  
export HDFS_DATANODE_USER="root"  
export HDFS_SECONDARYNAMENODE_USER="root"  
export YARN_RESOURCEMANAGER_USER="root"  
  
export YARN_NODEMANAGER_USER="root"
```

## [Hbase – 2.4]

```
# set to the root of your Java installation
export JAVA_HOME=/opt/jdk/

# Technically, the only required environment variable is JAVA_HOME.
# All others are optional. However, the defaults are probably not
# preferred. Many sites configure these options outside of Hadoop,
# such as in /etc/profile.d

export HDFS_NAMENODE_USER="root"
export HDFS_DATANODE_USER="root"
export HDFS_SECONDARYNAMENODE_USER="root"
export YARN_RESOURCEMANAGER_USER="root"
export YARN_NODEMANAGER_USER="root"

# The java implementation to use. By default, this environment
# variable is REQUIRED on ALL platforms except OS X!
export JAVA_HOME=/opt/jdk/

# Location of Hadoop. By default, Hadoop will attempt to determine
# this location based upon its execution path.
# export HADOOP_HOME=
```

Try the following command:

OSTech

## [Hbase – 2.4]

```
# cd /opt/hadoop
```

```
$ bin/hadoop
```

## Pseudo-Distributed Operation

Hadoop can also be run on a single-node in a pseudo-distributed mode where each Hadoop daemon runs in a separate Java process.

### Configuration

Use the following:

etc/hadoop/core-site.xml:

```
<configuration>
  <property>
    <name>fs.defaultFS</name>
    <value>hdfs://hmaster.master.com:9000</value>
  </property>
</configuration>
```

## [Hbase – 2.4]

:etc/hadoop/hdfs-site.xml:

```
<?xml version="1.0" encoding="UTF-8"?>
<?xml-stylesheet type="text/xsl" href="configuration.xsl"?>
<!--
    Licensed under the Apache License, Version 2.0 (the "License");
    you may not use this file except in compliance with the License.
    You may obtain a copy of the License at

        http://www.apache.org/licenses/LICENSE-2.0

    Unless required by applicable law or agreed to in writing, software
    distributed under the License is distributed on an "AS IS" BASIS,
    WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
    See the License for the specific language governing permissions and
    limitations under the License. See accompanying LICENSE file.
-->

<!-- Put site-specific property overrides in this file. -->

<configuration>
<property>
<name>dfs.namenode.name.dir</name>
<value>file:///opt/hdfs_hb0/namenode</value>
<description>NameNode directory for namespace and transaction logs storage.</description>
</property>
<property>
<name>dfs.datanode.data.dir</name>
<value>file:///opt/hdfs_hbo/datanode</value>
<description>DataNode directory</description>
</property>
<property>
<name>dfs.replication</name>
<value>2</value>
</property>
```

## [Hbase – 2.4]

```
<property>
<name>dfs.permissions</name>
<value>false</value>
</property>
<property>
<name>dfs.datanode.use.datanode.hostname</name>
<value>false</value>
</property>
<property>
<name>dfs.namenode.datanode.registration.ip-hostname-check</name>
<value>false</value>
</property>
<property>
<name>dfs.namenode.decommission.interval</name>
<value>30</value>
</property>
<property>
<name>dfs.replication</name>
<value>2</value>
</property>
<property>
<name>dfs.permissions</name>
<value>false</value>
</property>
<property>
<name>dfs.datanode.use.datanode.hostname</name>
<value>false</value>
</property>
<property>
<name>dfs.namenode.datanode.registration.ip-hostname-check</name>
<value>false</value>
</property>
<property>
<name>dfs.namenode.decommission.interval</name>
<value>30</value>
</property>
<property>
<name>dfs.client.datanode-restart.timeout</name>
<value>30</value>
</property>
```

## [Hbase – 2.4]

```
</configuration>
```

Setup passphraseless ssh

```
# yum install openssh*
```

change the password of root:

```
#passwd
```

Enter : root123

Start the ssh server.

```
# systemctl start sshd
```

Now check that you can ssh to the localhost without a passphrase:

```
$ ssh hmaster.master.com
```

If you cannot ssh to host without a passphrase, execute the following commands and try ssh again:

```
$ ssh-keygen -t rsa -P '' -f ~/.ssh/id_rsa  
$ cat ~/.ssh/id_rsa.pub >> ~/.ssh/authorized_keys  
$ chmod 0600 ~/.ssh/authorized_keys
```

## [Hbase – 2.4]

Create the necessary directory structure:

```
#mkdir -p /opt/hdfs_hb0/namenode
```

```
#mkdir -p /opt/hdfs_hb0/datanode
```

1. Format the filesystem:

```
2. $ bin/hdfs namenode -format
```

3. Start NameNode daemon and DataNode daemon:

```
4. $ sbin/start-dfs.sh
```

The hadoop daemon log output is written to the `$HADOOP_LOG_DIR` directory (defaults to `$HADOOP_HOME/logs`).

Verify the process : jps

```
[root@hmaster hadoop]# jps
4704 Jps
3876 DataNode
3753 NameNode
4076 SecondaryNameNode
[root@hmaster hadoop]#
```

## [Hbase – 2.4]

5. Browse the web interface for the NameNode; by default it is available at:
  - o NameNode - <http://localhost:9870/>
6. Make the HDFS directories required to execute MapReduce jobs:

```
7. $ bin/hdfs dfs -mkdir /user
```

```
$ bin/hdfs dfs -mkdir /user/root
```

```
[root@hmaster hadoop]# bin/hdfs dfs -mkdir /user
[root@hmaster hadoop]# bin/hdfs dfs -mkdir /user/root
[root@hmaster hadoop]# bin/hdfs dfs -ls /
Found 1 items
drwxr-xr-x  - root supergroup          0 2020-11-27 14:53 /user
[root@hmaster hadoop]# bin/hdfs dfs -ls -R /
drwxr-xr-x  - root supergroup          0 2020-11-27 14:53 /user
drwxr-xr-x  - root supergroup          0 2020-11-27 14:53 /user/root
[root@hmaster hadoop]#
```

## [Hbase – 2.4]

Following configuration are Optional – Not Required.

1. Configure parameters as follows:

etc/hadoop/mapred-site.xml:

```
<configuration>

    <property>
        <name>mapreduce.framework.name</name>
        <value>yarn</value>
    </property>
    <property>
        <name>mapreduce.application.classpath</name>
        <value>$HADOOP_MAPRED_HOME/share/hadoop/mapreduce/*:$HADOOP_MAPRED_HOME/share/hadoop/mapreduce/lib/*</value>
    </property>

</configuration>
```

## [Hbase – 2.4]

etc/hadoop/yarn-site.xml:

```
<configuration>
  <property>
    <name>yarn.nodemanager.aux-services</name>
    <value>mapreduce_shuffle</value>
  </property>
  <property>
    <name>yarn.nodemanager.env-whitelist</name>
    <value>JAVA_HOME,HADOOP_COMMON_HOME,HADOOP_HDFS_HOME,HADOOP_CONF_DIR,CLASSPATH_PREPEND_DISTCACHE,HADOOP_YARN_HOME,HADOOP_MAPRED_HOME</value>
  </property>
</configuration>
```

2. Start ResourceManager daemon and NodeManager daemon:

```
$ sbin/start-yarn.sh
```

```
[root@hmaster hadoop]# sbin/start-yarn.sh
Starting resourcemanager
Last login: Fri Nov 27 14:50:33 IST 2020 on pts/4
Starting nodemanagers
Last login: Fri Nov 27 14:59:20 IST 2020 on pts/4
```

3. Browse the web interface for the ResourceManager; by default it is available at:
  - o ResourceManager - <http://localhost:8088/>
4. Run a MapReduce job.
5. When you're done, stop the daemons with:

## [Hbase – 2.4]

```
$ sbin/stop-yarn.sh
```

----- Lab Ends Here -----

## **6. Using HDFS – 30 Minutes**

Most of your interaction with the system will be through commands supported by the bin/hdfs script. If you start a terminal and run this program with no arguments, it prints a help message. To try this, run the following command:

```
export PATH=$PATH:/opt/hadoop/bin:/opt/hadoop/sbin
```

```
$ hdfs
```

Enter:

```
$ hdfs dfs -ls /
```

This shows you the contents of the root directory in HDFS. There will be multiple entries, one of which is /user. Individual users have a “home” directory under. Folders may varies.

View the contents of the /hbase directory by entering the following:

```
$ hdfs dfs -ls /hbase
```

[Hbase – 2.4]

```
Subcommands may print help when invoked with parameters or when -h.
[root@hmaster opt]# hdfs dfs -ls /
Found 1 items
drwxr-xr-x  - root supergroup          0 2020-12-10 09:13 /hbase
[root@hmaster opt]# hdfs dfs -ls /hbase
Found 12 items
drwxr-xr-x  - root supergroup          0 2020-12-07 03:41 /hbase/.hbck
drwxr-xr-x  - root supergroup          0 2020-12-10 09:50 /hbase/.tmp
drwxr-xr-x  - root supergroup          0 2020-12-07 03:41 /hbase/MasterData
drwxr-xr-x  - root supergroup          0 2020-12-10 09:41 /hbase/WALs
drwxr-xr-x  - root supergroup          0 2020-12-08 07:55 /hbase/archive
drwxr-xr-x  - root supergroup          0 2020-12-07 03:41 /hbase/corrupt
drwxr-xr-x  - root supergroup          0 2020-12-07 03:41 /hbase/data
-rw-r--r--  3 root supergroup          42 2020-12-07 03:41 /hbase/hbase.id
-rw-r--r--  3 root supergroup          7 2020-12-07 03:41 /hbase/hbase.version
drwxr-xr-x  - root supergroup          0 2020-12-07 03:41 /hbase/mobdir
drwxr-xr-x  - root supergroup          0 2020-12-10 09:33 /hbase/oldWALs
drwx--x--x  - root supergroup          0 2020-12-07 03:41 /hbase/staging
[root@hmaster opt]#
```

Enter the following:

```
$ hdfs dfs -ls /data/training
```

There are no files, it would display an error message.

OSTech

## [Hbase – 2.4]

Note that the directory structure in HDFS has nothing to do with the directory structure of the local filesystem; they are completely separate namespaces.

### Step 2: Uploading Files

Besides browsing the existing filesystem, another important thing you can do with FsShell is to upload new data into HDFS.

1. Change directories to the directory containing the sample data we will be using in the course. Or select any folder that contains some files.

```
$ cd /opt/
```

```
[root@hmaster opt]#  
[root@hmaster opt]# pwd  
/opt  
[root@hmaster opt]# ls  
data  hadoop  hbase  jdk  shakespeare.tar  tmp  
[root@hmaster opt]#
```

If you perform a ‘regular’ ls command in this directory, you will see a few files, including one named shakespeare.tar.gz.

## [Hbase – 2.4]

Unzip shakespeare.tar.gz by running:

```
$ tar -xvf shakespeare.tar
```

This creates a directory named shakespeare/ containing several files on your local filesystem. If you have used this VM for a previous course, we recommend that you unzip the tar file again to ensure you have the file required for this lab.

Change to the shakespeare directory and view the files:

```
$ cd shakespeare  
$ ls
```

You will see a few files, one of which is named glossary.

Insert this file into HDFS:

```
$ hdfs dfs -mkdir -p /user/training/glossary  
$ hdfs dfs -put glossary /user/training/glossary
```

This copies the local glossary file into the remote HDFS directory named /user/training/

List the contents of your HDFS home directory now:

OSTech

[Hbase – 2.4]

```
$ hdfs dfs -ls /user/training
```

You should see an entry for the glossary file.

Enter:

```
$ hdfs dfs -tail /user/training/glossary/glossary
```

```
XANTHIPPE      Socrate's scolding wife

YARE      ready, being understood
YARELY    readily
YAW       out of control
Y-CLAD    clad
Y-CLEPED   called, named
YEARN     to grieve, vex

YELLOWNESS   jealousy
YELLOWS     a disease of horses
YEOMAN      a sheriff's officer
YIELD      to reward
            To report
YOND       and yonder
YOUNKER    tyro

ZANY      a clown, gull
[END]
[root@hmaster shakespeare]# hdfs dfs -tail /user/training/glossary/glossary
```

OSTech

## [Hbase – 2.4]

This prints the last 1KB of the glossary to your terminal. This command is handy for viewing the output of MapReduce programs. Very often, an individual output file of a MapReduce program is very large, making it inconvenient to view the entire file in the terminal.

### Other Commands

There are several other commands associated with the FsShell subsystem which let you perform most common filesystem manipulations: rm, rm -r (recursive rm), mv, cp, mkdir, etc.

1. In the terminal window, enter:

```
$ hdfs dfs
```

You see a help message describing all the commands associated with this subsystem.

### Bonus Exercise: Advanced FsShell Commands

1. In the terminal window, enter:

```
$ hdfs dfs -help
```

This will give a more complete description of each FsShell command and the arguments for it.

2. Enter:

OSTech

## [Hbase – 2.4]

```
$ hdfs dfs -cat /user/training/glossary/glossary | tail -n 50
```

This prints the last 50 lines of the glossary to your terminal. This command is handy for viewing the output of MapReduce programs. Very often, an individual output file of a MapReduce program is very large, making it inconvenient to view the entire file in the terminal. For this reason, it's often a good idea to pipe the output of the fs -cat command into head, tail, more, or less.

----- Lab Ends Here -----

## 7. Pseudo Distributed Hbase Cluster – 60 Minutes

This lab depends on the Lab – Configure Hadoop Cluster.

Each HBase daemon (HMaster, HRegionServer, and ZooKeeper) runs as a separate process.

By default, unless you configure the `hbase.rootdir` property , your data is still stored in `/tmp/`.

Stop HBase if it is running

Configure HBase.

Edit the conf/`hbase-site.xml` configuration. First, add the following property which directs HBase to run in distributed mode, with one JVM instance per daemon.

```
<property>
  <name>hbase.cluster.distributed</name>
  <value>true</value>
</property>
<property>
  <name>hbase.rootdir</name>
  <value>hdfs://hmaster.master.com:9000/hbase</value>
</property>
```

Skip the following installation if Already done before else execute the following commands to install the ssh server.

```
#yum install openssh-server openssh-clients
```

```
#systemctl start sshd
```

OSTech

## [Hbase – 2.4]

Start HBase.

Use the *bin/start-hbase.sh* command to start HBase. If your system is configured correctly, the *jps* command should show the HMaster and HRegionServer processes running.

```
[root@hmaster hbase]# bin/start-hbase.sh
127.0.0.1: running zookeeper, logging to /opt/hbase/bin/..../logs/hbase-root-zookeeper-hmaster.master.com.out
running master, logging to /opt/hbase/bin/..../logs/hbase--master-hmaster.master.com.out
hmaster.master.com: running regionserver, logging to /opt/hbase/bin/..../logs/hbase-root-regionserver-hmaster.master
.com.out
[root@hmaster hbase]# jps
3876 DataNode
10293 HMaster
3753 NameNode
10538 Jps
10426 HRegionServer
10188 HQuorumPeer
4076 SecondaryNameNode
[root@hmaster hbase]#
```

Check the HBase directory

Open a terminal, change directory /opt/hadoop

```
# hdfs dfs -ls /hbase
```

OSTech

## [Hbase – 2.4]

```
[root@hmaster hadoop]# bin/hadoop fs -ls /hbase
Found 12 items
drwxr-xr-x  - root supergroup          0 2020-11-27 15:44 /hbase/.hbck
drwxr-xr-x  - root supergroup          0 2020-11-27 15:44 /hbase/.tmp
drwxr-xr-x  - root supergroup          0 2020-11-27 15:44 /hbase/MasterData
drwxr-xr-x  - root supergroup          0 2020-11-27 15:44 /hbase/WALs
drwxr-xr-x  - root supergroup          0 2020-11-27 15:44 /hbase/archive
drwxr-xr-x  - root supergroup          0 2020-11-27 15:44 /hbase/corrupt
drwxr-xr-x  - root supergroup          0 2020-11-27 15:44 /hbase/data
-rw-r--r--  3 root supergroup          42 2020-11-27 15:44 /hbase/hbase.id
-rw-r--r--  3 root supergroup          7 2020-11-27 15:44 /hbase/hbase.version
drwxr-xr-x  - root supergroup          0 2020-11-27 15:44 /hbase/mobdir
drwxr-xr-x  - root supergroup          0 2020-11-27 15:44 /hbase/oldWALs
drwx--x--x  - root supergroup          0 2020-11-27 15:44 /hbase/staging
[root@hmaster hadoop]#
```

Create a table and populate it with data.

```
#bin/hbase shell
create 'test', 'cf'
list 'test'
put 'test', 'row1', 'cf:a', 'value1'
put 'test', 'row2', 'cf:b', 'value2'
put 'test', 'row3', 'cf:c', 'value3'
scan 'test'
```

OSTech

## [Hbase – 2.4]

```
Version 2.3.3, r3e4bf4bee3a08b25591b9c22fea0518686a7e834, Wed Oct 28 06:36:25 UTC 2020
Took 0.0042 seconds
hbase(main):001:0> create 'test', 'cf'
Created table test
Took 2.7289 seconds
=> Hbase::Table - test
hbase(main):002:0> list 'test'
TABLE
test
1 row(s)
Took 0.0771 seconds
=> ["test"]
hbase(main):003:0> put 'test', 'row1', 'cf:a', 'value1'
Took 0.5316 seconds
hbase(main):004:0> put 'test', 'row2', 'cf:b', 'value2'
Took 0.0135 seconds
hbase(main):005:0> put 'test', 'row3', 'cf:c', 'value3'
Took 0.0138 seconds
hbase(main):006:0> scan 'test'
ROW                                COLUMN+CELL
row1                               column=cf:a, timestamp=2020-11-27T16:00:27.341, value=value1
row2                               column=cf:b, timestamp=2020-11-27T16:00:32.395, value=value2
row3                               column=cf:c, timestamp=2020-11-27T16:00:37.782, value=value3
3 row(s)
Took 0.1571 seconds
hbase(main):007:0>
```

----- Lab Ends Here -----

*Procedure: Stop HBase*

1. In the same way that the *bin/start-hbase.sh* script is provided to conveniently start all HBase daemons, the *bin/stop-hbase.sh* script stops them.

```
2. $ ./bin/stop-hbase.sh  
3. stopping hbase.....  
$
```

4. After issuing the command, it can take several minutes for the processes to shut down. Use the *jps* to be sure that the HMaster and HRegionServer processes are shut down.

The above has shown you how to start and stop a standalone instance of HBase. In the next sections we give a quick overview of other modes of hbase deploy.

<https://hbase.apache.org/book.html#quickstart>

## 8. Using the Developer JAVA API - 60 Minutes

In this exercise you will use the HBase Java API to programmatically perform CRUD operations using - Put, Get, and Scan methods.

Requirements:

- Mention the zookeeper details in hbase-site.xml.
- Start Hbase Pseudo cluster.
- Use eclipse to write code.

Use Case: Create a table and insert records in the table.

Start the eclipse and create a maven project with the following project details. Use simple maven project type.

```
<groupId>com.ostechn&lt;/groupId>
<artifactId>LearningHbase</artifactId>
<version>0.0.1-SNAPSHOT</version>
```

Update the pom.xml with the following details. All jars dependencies are being mention in this file.

[Hbase – 2.4]

```
<project xmlns="http://maven.apache.org/POM/4.0.0"
  xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
  xsi:schemaLocation="http://maven.apache.org/POM/4.0.0
https://maven.apache.org/xsd/maven-4.0.0.xsd">
  <modelVersion>4.0.0</modelVersion>
  <groupId>com.ostechn&lt;/groupId>
  <artifactId>LearningHbase</artifactId>
  <version>0.0.1-SNAPSHOT</version>
  <properties>
    <maven.compiler.source>1.8</maven.compiler.source>
    <maven.compiler.target>1.8</maven.compiler.target>
    <project.build.sourceEncoding>UTF-
8</project.build.sourceEncoding>
    <junit.version>4.8.1</junit.version>
    <hadoop.version>2.7.2</hadoop.version>
    <hbase.version>2.3.3</hbase.version>
  </properties>
  <dependencies>

    <!-- Hadoop -->
    <dependency>
      <groupId>org.apache.hadoop</groupId>
      <artifactId>hadoop-hdfs</artifactId>
      <version>${hadoop.version}</version>
```

OSTech

[Hbase – 2.4]

```
<exclusions>
    <exclusion>
        <groupId>javax.servlet</groupId>
        <artifactId>*</artifactId>
    </exclusion>
</exclusions>
</dependency>
<dependency>
    <groupId>org.apache.hadoop</groupId>
    <artifactId>hadoop-client</artifactId>
    <version>${hadoop.version}</version>
    <exclusions>
        <exclusion>
            <groupId>javax.servlet</groupId>
            <artifactId>*</artifactId>
        </exclusion>
    </exclusions>
</dependency>

<!-- hbase version -->
<dependency>
    <groupId>org.apache.hbase</groupId>
    <artifactId>hbase-client</artifactId>
    <version>${hbase.version}</version>
```

OSTech

[Hbase – 2.4]

```
</dependency>
```

```
</dependencies>  
</project>
```

Create a package with the following class in it.

```
package com.ostechn.hbase;  
  
import java.io.IOException;  
  
import org.apache.hadoop.conf.Configuration;  
import org.apache.hadoop.hbase.HBaseConfiguration;  
import org.apache.hadoop.hbase.TableName;  
import org.apache.hadoop.hbase.client.Admin;  
import org.apache.hadoop.hbase.client.ColumnFamilyDescriptorBuilder;  
import org.apache.hadoop.hbase.client.Connection;  
import org.apache.hadoop.hbase.client.ConnectionFactory;  
import org.apache.hadoop.hbase.client.Put;  
import org.apache.hadoop.hbase.client.Table;  
import org.apache.hadoop.hbase.client.TableDescriptor;  
import org.apache.hadoop.hbase.client.TableDescriptorBuilder;
```

OSTech

[Hbase – 2.4]

```
import org.apache.hadoop.hbase.util.Bytes;

public class UpdateMovies {

    public static void main(String[] args) throws IOException {
        // TODO Auto-generated method stub
        Configuration config = HBaseConfiguration.create();
        String path = UpdateMovies.class
            .getClassLoader()
            .getResource("hbase-site.xml")
            .getPath();

        Connection connection =
ConnectionFactory.createConnection(config);
        try {
            Admin admin = connection.getAdmin() ;

            System.out.println(">>>" + admin);
            createTable(admin);
            putRow(connection.getTable(TABLE_NAME));
            System.out.println(">>> Successfully Created the table and
inserted a row " + admin);

        }catch (Exception e) {
```

OSTech

[Hbase – 2.4]

```
        e.printStackTrace();
    }

}

private static final TableName TABLE_NAME =
TableName.valueOf("mymovietbl");
private static final byte[] CF_NAME = Bytes.toBytes("movie");
private static final byte[] QUALIFIER = Bytes.toBytes("mname");
private static final byte[] ROW_ID = Bytes.toBytes("row01");

public static void createTable(final Admin admin) throws
IOException {
    if (!admin.tableExists(TABLE_NAME)) {
        TableDescriptor desc =
TableDescriptorBuilder.newBuilder(TABLE_NAME)
.setColumnFamily(ColumnFamilyDescriptorBuilder.of(CF_NAME))
.build();
        admin.createTable(desc);
    }
}
```

OSTech

[Hbase – 2.4]

```
public static void putRow(final Table table) throws IOException
{
    table.put(new Put(ROW_ID).addColumn(CF_NAME, QUALIFIER,
Bytes.toBytes("Hello, World!")));
}
```

In the above class, we have mention `hbase-site.xml` in the configuration to fetch the hbase cluster details for creating a table.

Copy the following `hbase-site.xml` content in the the following location:

## [Hbase – 2.4]

The screenshot shows the Eclipse IDE interface with the title "eclipse-workspace - LearningHbase/src/main/resources/hbase-site.xml - Eclipse IDE". The left pane displays the "Package Explorer" with project files like CFG, couchbase, HelloCouchbase, and LearningHbase. The "src/main/java" folder contains com.ostechn.hbase subfolder with LearningHbase.java, MyLittleHBaseClient.java, and UpdateMovies.java. The "src/main/resources" folder contains hbase-site.xml. A red arrow points from the "src/main/resources" folder to the hbase-site.xml file in the central editor area. The right pane shows the "Console" tab with log output:

```
<terminated> UpdateMovies [Java Application] /Library/Java/JavaVirtualMachines/jdk1.8.0_271.jdk/Contents/Home/bin/java (08-Dec-2018) log4j:WARN No appenders could be found for logger (org.apache.hadoop.util.Shell). log4j:WARN Please initialize the log4j system properly. log4j:WARN See http://logging.apache.org/log4j/1.2/faq.html#noconfig for more info. >>>org.apache.hadoop.hbase.client.HBaseAdmin@55740540>>> Successfully Created the table and inserted
```

[Hbase – 2.4]

```
<?xml version="1.0"?>
<?xml-stylesheet type="text/xsl" href="configuration.xsl"?>
<configuration>
  <property>
    <name>hbase.zookeeper.quorum</name>
    <value>localhost</value>
  </property>
  <property>
    <name>hbase.zookeeper.property.clientPort</name>
    <value>2181</value>
  </property>
</configuration>
```

Execute the main program.

You can verify the result from the hbase shell. Start a hbase console.

#hbase shell

Determine the table in the shell – mymovietbl

```
#list
#scan 'mymovietbl'
```

OSTech

## [Hbase – 2.4]

```
hbase(main):007:0* list
TABLE
mymovietbl
test
2 row(s)
Took 0.0565 seconds
=> ["mymovietbl", "test"]
hbase(main):008:0> scan 'mymovietbl'
ROW                                     COLUMN+CELL
row01                                  column=movie:mname, timestamp=2020-12-08T08:11:19.876, value=Hello, World!
1 row(s)
Took 0.1123 seconds
hbase(main):009:0> █
```

As shown above, the table; mymovietbl was created and a row with row01 was inserted.

In the next program, we will perform the followings:

- Insert 3 movies record using Put.
- Fetch a record using – Get (Row Key - nungshi)
- Fetch all records using scanner.

[Hbase – 2.4]

Create the following class.

```
package co.ostech.hbase;

import java.io.IOException;
import java.util.ArrayList;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.hbase.CompareOperator;
import org.apache.hadoop.hbase.HBaseConfiguration;
import org.apache.hadoop.hbase.TableName;
import org.apache.hadoop.hbase.client.Connection;
import org.apache.hadoop.hbase.client.ConnectionFactory;
import org.apache.hadoop.hbase.client.Get;
import org.apache.hadoop.hbase.client.Put;
import org.apache.hadoop.hbase.client.Result;
import org.apache.hadoop.hbase.client.ResultScanner;
import org.apache.hadoop.hbase.client.Scan;
import org.apache.hadoop.hbase.client.Table;
import org.apache.hadoop.hbase.filter.Filter;
import org.apache.hadoop.hbase.filter.FilterList;
import org.apache.hadoop.hbase.filter.SingleColumnValueFilter;
```

OSTech

[Hbase – 2.4]

```
import org.apache.hadoop.hbase.util.Bytes;

public class MyLittleHBaseClient {
    private static String table_name = "mymovietbl";
    public static void main(String[] args) throws IOException {
        Configuration config = HBaseConfiguration.create();
        String path =
UpdateMovies.class.getClassLoader().getResource("hbase-
site.xml").getPath();
        config.addResource(new Path(path));
        Connection connection =
ConnectionFactory.createConnection(config);

        try {
            Table table =
connection.getTable(TableName.valueOf(table_name));
            UpdateInfo(table);
            //findMoviesByLanguage(table);
            //findMoviesByTwoLanguage(table);
        } finally {
            connection.close();
        }
    }
}
```

OSTech

[Hbase – 2.4]

```
public static void UpdateInfo(Table table) throws IOException {  
    try {  
        Put p = new Put(Bytes.toBytes("braveheart"));  
        p.addColumn(Bytes.toBytes("movie"), Bytes.toBytes("mname"), Bytes.toBytes("Braveheart"));  
        p.addColumn(Bytes.toBytes("movie"), Bytes.toBytes("language"), Bytes.toBytes("English"));  
        table.put(p);  
        Put p1 = new Put(Bytes.toBytes("kasam"));  
        p1.addColumn(Bytes.toBytes("movie"), Bytes.toBytes("mname"), Bytes.toBytes("Kasam"));  
        p1.addColumn(Bytes.toBytes("movie"), Bytes.toBytes("language"), Bytes.toBytes("Hindi"));  
        table.put(p1);  
    } catch (IOException e) {  
        e.printStackTrace();  
    }  
}
```

OSTech

[Hbase – 2.4]

```
Put p2 = new Put(Bytes.toBytes("nungshi"));

    p2.addColumn(Bytes.toBytes("movie"), Bytes.toBytes("mname"), Bytes.toB
ytes("Nungshi"));

    p2.addColumn(Bytes.toBytes("movie"), Bytes.toBytes("language"), Bytes.
toBytes("Manipuri"));

    table.put(p2);

    Get g = new Get(Bytes.toBytes("nungshi"));
    Result r = table.get(g);
    byte[] value = r.getValue(Bytes.toBytes("movie"),
Bytes.toBytes("mname"));

        // If we convert the value bytes, we should get back 'Some
Value', the
        // value we inserted at this location.
        String valueStr = Bytes.toString(value);
        System.out.println("Get Information for Row Key - nungshi " +
valueStr);
```

OSTech

[Hbase – 2.4]

```
// Sometimes, you won't know the row you're looking for. In
this case, you
    // use a Scanner. This will give you cursor-like interface to
the contents
        // of the table. To set up a Scanner, do like you did above
making a Put
            // and a Get, create a Scan. Adorn it with column names, etc.
Scan s = new Scan();
s.addColumn(Bytes.toBytes("movie"), Bytes.toBytes("mname"));
ResultScanner scanner = table.getScanner(s);
try {
    // Scanners return Result instances.
    // Now, for the actual iteration. One way is to use a
while loop like so:
    for (Result rr = scanner.next(); rr != null; rr =
scanner.next()) {
        // print out the row we found and the columns we were
looking for
            System.out.println("Found row: " + rr);
    }
}
// The other approach is to use a foreach loop. Scanners
are iterable!
// for (Result rr : scanner) {
```

OSTech

[Hbase – 2.4]

```
        // System.out.println("Found row: " + rr);
        // }
    } finally {
    // Make sure you close your scanners when you are done!
    // Thats why we have it inside a try/finally clause
    scanner.close();
}

// Close your table and cluster connection.
} finally {
if (table != null)
    table.close();
}
}

public static void findMoviesByLanguage(Table table) throws
IOException {
    byte[] FAMILY_BYTES = Bytes.toBytes("movie"); // Column Family
Name
    byte[] COLUMN_BYTES = Bytes.toBytes("language"); // Column
    SingleColumnValueFilter f1 = new
SingleColumnValueFilter(FAMILY_BYTES,COLUMN_BYTES,
CompareOperator.EQUAL, Bytes.toBytes("Manipuri"));
OSTech
```

[Hbase – 2.4]

```
f1.setFilterIfMissing(true);

ArrayList<Filter> filters = new ArrayList<Filter>();
filters.add(f1);

FilterList filterList = new
FilterList(FilterList.Operator.MUST_PASS_ONE, filters);

Scan s = new Scan();
s.setFilter(filterList);
s.addColumn(Bytes.toBytes("movie"), Bytes.toBytes("mname"));
s.addColumn(Bytes.toBytes("movie"), Bytes.toBytes("language"));

ResultScanner scanner = table.getScanner(s);
for (Result rr = scanner.next(); rr != null; rr =
scanner.next()) {
    // print out the row we found and the columns we were looking
for
    System.out.println("Language: " +
        Bytes.toString(rr.getValue(Bytes.toBytes("movie"),
Bytes.toBytes("language")))
        + "\t" +
        Bytes.toString(rr.getValue(Bytes.toBytes("movie"),
Bytes.toBytes("mname"))));
}
```

OSTech

[Hbase – 2.4]

}

}

```
public static void findMoviesByTwoLanguage(Table table) throws  
IOException {  
  
    byte[] FAMILY_BYTES = Bytes.toBytes("movie");  
    byte[] COLUMN_BYTES = Bytes.toBytes("language");  
    SingleColumnValueFilter f1 = new  
SingleColumnValueFilter(FAMILY_BYTES,COLUMN_BYTES,  
CompareOperator.EQUAL, Bytes.toBytes("Manipuri"));  
    f1.setFilterIfMissing(true);  
  
    SingleColumnValueFilter f2 = new  
SingleColumnValueFilter(FAMILY_BYTES,COLUMN_BYTES,  
CompareOperator.EQUAL, Bytes.toBytes("English"));  
    f2.setFilterIfMissing(true);  
  
    ArrayList<Filter> filters = new ArrayList<Filter>();  
    filters.add(f1);  
    filters.add(f2);
```

OSTech

[Hbase – 2.4]

```
FilterList filterList = new  
FilterList(FilterList.Operator.MUST_PASS_ONE, filters);  
  
Scan s = new Scan();  
s.setFilter(filterList);  
s.addColumn(Bytes.toBytes("movie"), Bytes.toBytes("mname"));  
s.addColumn(Bytes.toBytes("movie"), Bytes.toBytes("language"));  
  
ResultScanner scanner = table.getScanner(s);  
for (Result rr = scanner.next(); rr != null; rr =  
scanner.next()) {  
    // print out the row we found and the columns we were looking  
    for  
        System.out.println("Language: " +  
            Bytes.toString(rr.getValue(Bytes.toBytes("movie"),  
Bytes.toBytes("language")))  
            + "\t" +  
            Bytes.toString(rr.getValue(Bytes.toBytes("movie"),  
Bytes.toBytes("mname"))));  
}
```

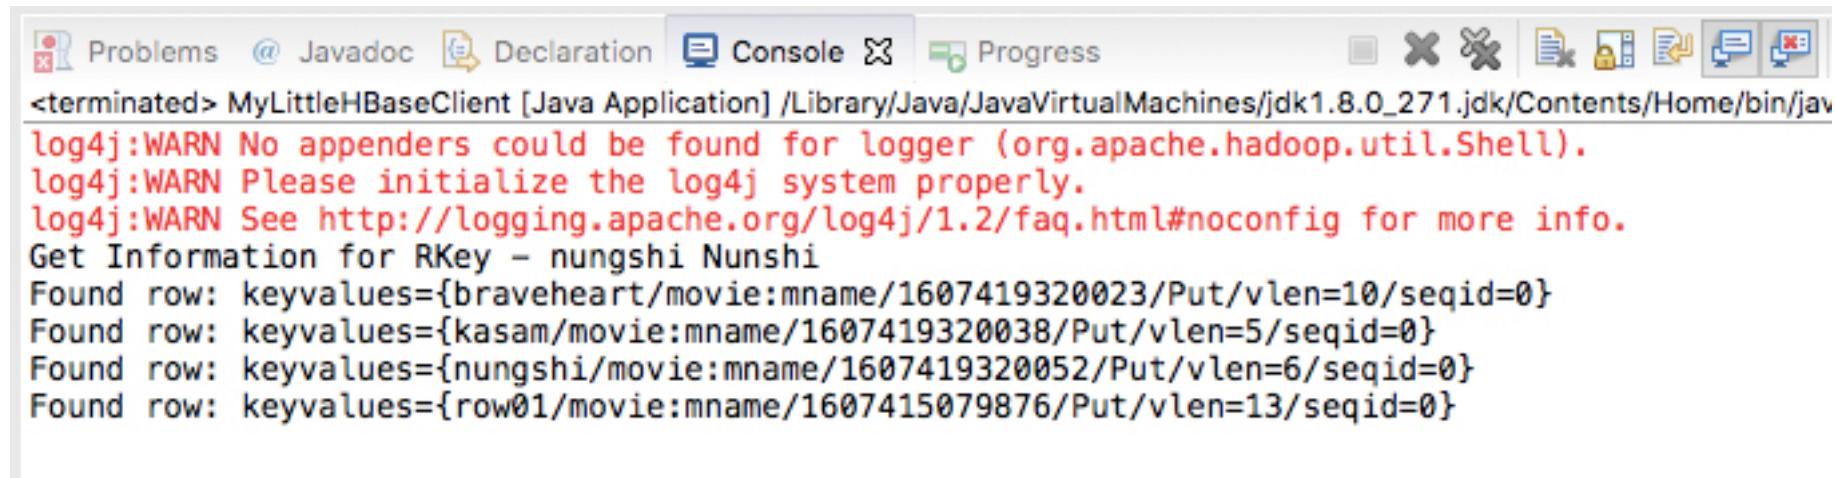
OSTech

[Hbase – 2.4]

}

}

Execute the program:



The screenshot shows the Eclipse IDE interface with the 'Console' tab selected. The output window displays the following text:

```
<terminated> MyLittleHBaseClient [Java Application] /Library/Java/JavaVirtualMachines/jdk1.8.0_271.jdk/Contents/Home/bin/java
log4j:WARN No appenders could be found for logger (org.apache.hadoop.util.Shell).
log4j:WARN Please initialize the log4j system properly.
log4j:WARN See http://logging.apache.org/log4j/1.2/faq.html#noconfig for more info.
Get Information for RKey - nungshi Nunshi
Found row: keyvalues={braveheart/movie:mname/1607419320023/Put/vlen=10/seqid=0}
Found row: keyvalues={kasam/movie:mname/1607419320038/Put/vlen=5/seqid=0}
Found row: keyvalues={nungshi/movie:mname/1607419320052/Put/vlen=6/seqid=0}
Found row: keyvalues={row01/movie:mname/1607415079876/Put/vlen=13/seqid=0}
```

As shown above, the first row is the result of the record being fetch using key – nungshi. The last four rows are fetch using the scanner.

## [Hbase – 2.4]

Case 1:

Comment and uncomment as shown below in the main method.  
It will apply a single filter to fetch the record with the language of type ‘Manipuri’

```
J MyLittleHBaseClient.java 
29     config.addResource(new Path(path));
30     Connection connection = ConnectionFactory.createConnection(config);
31
32     try {
33
34         Table table = connection.getTable(TableName.valueOf(table_name));
35         //UpdateInfo(table);
36         ✓findMoviesByLanguage(table);
37         //findMoviesByTwoLanguage(table);
38     } finally {
39         connection.close();
40     }
41 }
42
43⊕ public static void UpdateInfo(Table table) throws IOException {
44 }
```

Execute the program.

As shown below. Its filter by the Manipuri.

OSTech

## [Hbase – 2.4]

```
1  ~~~e(Table table) throws IOException {  
2  107     es("movie"); // Column Family Name  
3  108     es("language"); // Column  
4  109     ingleColumnValueFilter(FAMILY_BYTES,COLUMN_BYTES, CompareOperator.EQUAL, Bytes.toBytes("Manipuri"));  
5  110  
6  111  
7  112  
8  113  
9  114     rayList<Filter>();  
10 115  
11 116
```



Next apply Two filters i.e English and Manipuri Language.

Comment as shown below:

## [Hbase – 2.4]

```
31
32     try {
33
34         Table table = connection.getTable(TableName.valueOf(table_name));
35         //UpdateInfo(table);
36         // findMoviesByLanguage(table);
37         ✓findMoviesByTwoLanguage(table);
38     } finally {
39         connection.close();
40     }
41 }
42
43✉ public static void UpdateInfo(Table table) throws IOException {
44
45     try {
```

Execute the program.

## [Hbase – 2.4]

```
140 singleColumnValueFilter(FAMILY_BYTES,COLUMN_BYTES, CompareOperator.EQUAL, Bytes.toBytes("Manipuri"));
141
142
143 singleColumnValueFilter(FAMILY_BYTES,COLUMN_BYTES, CompareOperator.EQUAL, Bytes.toBytes("English"));
144
145
146 rayList<Filter>();
147
148
```

The screenshot shows a Java application window with a console tab. The console output displays two rows of data from a HBase table. The first row has 'Language: English' under the first column and 'Braveheart' under the second column. The second row has 'Language: Manipuri' under the first column and 'Nungshi' under the second column. Both 'Language' entries are crossed out with a red mark.

```
Progress Console X
<terminated> MyLittleHBaseClient (1) [Java Application] /Library/Java/JavaVirtualMachines/jdk1.8.0_271.jdk/Contents/Home/bin/java (06-Jun-2022, 8:45:18 AM – 8:45:30 AM)
Language: English      Braveheart
Language: Manipuri    Nungshi
```

Task:

Try adding a new record and execute scanner to fetch that record.

----- Lab Ends Here -----

OSTech

[Hbase – 2.4]

## 9. HBase Filters – 30 Minutes

In this exercise you will use the HBase Java API to programmatically scan rows in a hbase table and apply Filter on it.

With HBase scan and filters, it verify every row in the mymovietbl table and find those rows that match the following criteria:

1. Language is Manipuri
2. Language is Manipuri or English

It displays the movie name, and language for those rows that match the above criterias.

Use the Java Project created in the Lab- Using the Developer API

1 ) Add the following method in the class file, **MyLittleHBaseClient . java** and invoke from the main method. Comment all others method invocation from the main method. – Filter by One Column Value Only.

```
public static void findMoviesByLanguage(Table table) throws IOException {  
    byte[] FAMILY_BYTES = Bytes.toBytes("movie"); // Column Family  
    Name  
    byte[] COLUMN_BYTES = Bytes.toBytes("language"); // Column  
    OSTech
```

[Hbase – 2.4]

```
SingleColumnValueFilter f1 = new
SingleColumnValueFilter(FAMILY_BYTES,COLUMN_BYTES,
CompareOperator.EQUAL, Bytes.toBytes("Manipuri"));
f1.setFilterIfMissing(true);

ArrayList<Filter> filters = new ArrayList<Filter>();
filters.add(f1);

FilterList filterList = new
FilterList(FilterList.Operator.MUST_PASS_ONE, filters);

Scan s = new Scan();
s.setFilter(filterList);
s.addColumn(Bytes.toBytes("movie"), Bytes.toBytes("mname"));
s.addColumn(Bytes.toBytes("movie"), Bytes.toBytes("language"));

ResultScanner scanner = table.getScanner(s);
for (Result rr = scanner.next(); rr != null; rr =
scanner.next()) {
    // print out the row we found and the columns we were looking
for
    System.out.println("Language: " +
        Bytes.toString(rr.getValue(Bytes.toBytes("movie"),
Bytes.toBytes("language"))))
```

OSTech

[Hbase – 2.4]

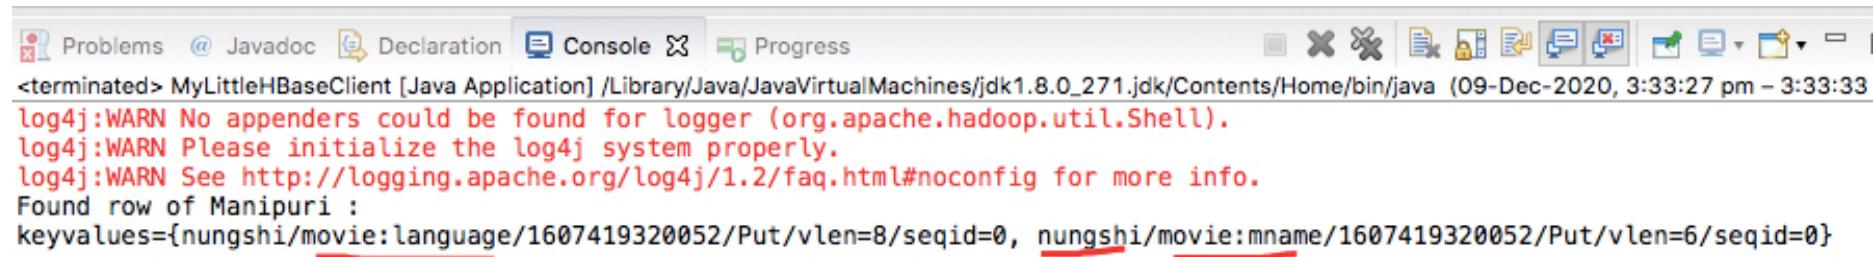
```
        + "\t" +
Bytes.toString(rr.getValue(Bytes.toBytes("movie"),
Bytes.toBytes("mname"))));
}

}
-- -----
23
24 public class MyLittleHBaseClient {
25     private static String table_name = "mymovietbl";
26     public static void main(String[] args) throws IOException {
27         Configuration config = HBaseConfiguration.create();
28         String path = UpdateMovies.class.getClassLoader().getResource("hbase-site.xml").getPath();
29         config.addResource(new Path(path));
30         Connection connection = ConnectionFactory.createConnection(config);
31
32     try {
33
34         Table table = connection.getTable(TableName.valueOf(table_name));
35         //UpdateInfo(table);
36         findMoviesByLanguage(table); ]
37         //findMoviesByTwoLanguage(table); |
38     } finally {
39         connection.close();
40     }
41 }
42
```

OSTech

## [Hbase – 2.4]

Execute the class.



```
Problems @ Javadoc Declaration Console Progress
<terminated> MyLittleHBaseClient [Java Application] /Library/Java/JavaVirtualMachines/jdk1.8.0_271.jdk/Contents/Home/bin/java (09-Dec-2020, 3:33:27 pm - 3:33:33)
log4j:WARN No appenders could be found for logger (org.apache.hadoop.util.Shell).
log4j:WARN Please initialize the log4j system properly.
log4j:WARN See http://logging.apache.org/log4j/1.2/faq.html#noconfig for more info.
Found row of Manipuri :
keyvalues={nungshi/movie:language/1607419320052/Put/vlen=8/seqid=0, nungshi/movie:mname/1607419320052/Put/vlen=6/seqid=0}
```

You will get only the Manipuri movie as shown above.

Now let us filter the English and Manipuri movies i.e applying multiples filters and will select specific columns only.

Add the following method.

```
public static void findMoviesByTwoLanguage(Table table) throws
IOException {

    byte[] FAMILY_BYTES = Bytes.toBytes("movie");
    byte[] COLUMN_BYTES = Bytes.toBytes("language");
    SingleColumnValueFilter f1 = new
    SingleColumnValueFilter(FAMILY_BYTES,COLUMN_BYTES,
    CompareOperator.EQUAL, Bytes.toBytes("Manipuri"));
```

OSTech

[Hbase – 2.4]

```
f1.setFilterIfMissing(true);

SingleColumnValueFilter f2 = new
SingleColumnValueFilter(FAMILY_BYTS, COLUMN_BYTS,
CompareOperator.EQUAL, Bytes.toBytes("English"));
f2.setFilterIfMissing(true);

ArrayList<Filter> filters = new ArrayList<Filter>();
filters.add(f1);
filters.add(f2);

FilterList filterList = new
FilterList(FilterList.Operator.MUST_PASS_ONE, filters);

Scan s = new Scan();
s.setFilter(filterList);
s.addColumn(Bytes.toBytes("movie"), Bytes.toBytes("mname"));
s.addColumn(Bytes.toBytes("movie"), Bytes.toBytes("language"));

ResultScanner scanner = table.getScanner(s);
for (Result rr = scanner.next(); rr != null; rr =
scanner.next()) {
    // print out the row we found and the columns we were looking
for
```

OSTech

[Hbase – 2.4]

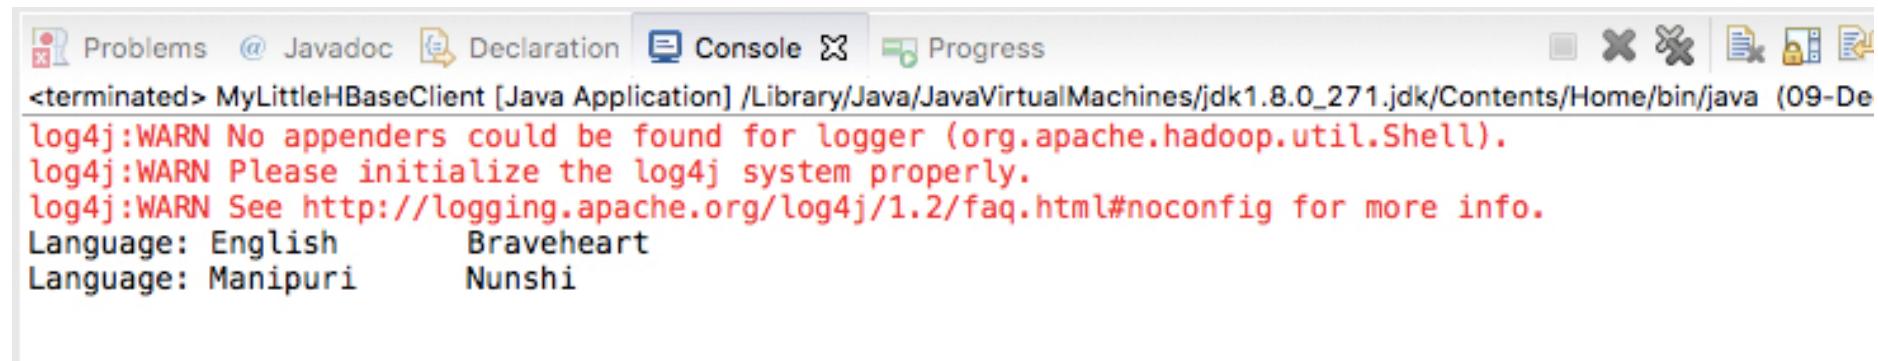
```
System.out.println("Language: " +
    Bytes.toString(rr.getValue(Bytes.toBytes("movie"),
Bytes.toBytes("language"))))
    + "\t" +
Bytes.toString(rr.getValue(Bytes.toBytes("movie"),
Bytes.toBytes("mname")))) );  
}  
}
```

Now invoke the above method by commenting earlier methods.

## [Hbase – 2.4]

```
~ 4 public class MyLittleHBaseClient {
5     private static String table_name = "mymovietbl";
6     public static void main(String[] args) throws IOException {
7         Configuration config = HBaseConfiguration.create();
8         String path = UpdateMovies.class.getClassLoader().getResource("hbase-site.xml").getPath();
9         config.addResource(new Path(path));
0         Connection connection = ConnectionFactory.createConnection(config);
1
2     try {
3
4         Table table = connection.getTable(TableName.valueOf(table_name));
5         //UpdateInfo(table);
6         //findMoviesByLanguage(table);
7         findMoviesByTwoLanguage(table);
8     } finally {
9         connection.close();
0     }
1 }
2 }
```

Execute the main class:



The screenshot shows the Eclipse IDE interface with the 'Console' tab selected. The output window displays the following text:

```
<terminated> MyLittleHBaseClient [Java Application] /Library/Java/JavaVirtualMachines/jdk1.8.0_271.jdk/Contents/Home/bin/java (09-De
log4j:WARN No appenders could be found for logger (org.apache.hadoop.util.Shell).
log4j:WARN Please initialize the log4j system properly.
log4j:WARN See http://logging.apache.org/log4j/1.2/faq.html#noconfig for more info.
Language: English      Braveheart
Language: Manipuri    Nunshi
```

[Hbase – 2.4]

As shown above, two filters were applied and only the movies of English and Manipuri were displayed.

----- Lab Ends here -----

## [Hbase – 2.4]

### 10. Built-in TSV Bulk Loader – 60 Minutes

ImportTsv is a utility that will load data in TSV format into HBase. It has two distinct usages:

- loading data from TSV format in HDFS into HBase via Puts,
- and preparing StoreFiles to be loaded via the `complexbulkload`.

In this lab, we are going to perform the second option.

Start HDFS cluster.

Start hbase cluster

HBase ships with a MR job that can read a delimiter-separated values file and output directly into an HBase table or create HFiles for bulk loading. Here we are going to:

1. Get the sample data and upload it to HDFS. - `word_count.csv`
2. Run the ImportTsv job to transform the file into multiple HFiles according to a pre-configured table.
3. Prepare and load the files in HBase.

Now, upload the file to HDFS:

Create /data directory to upload the `word_count.csv` file to it.

```
#hdfs dfs -mkdir /data
```

```
#hdfs dfs -put word_count.csv /data
```

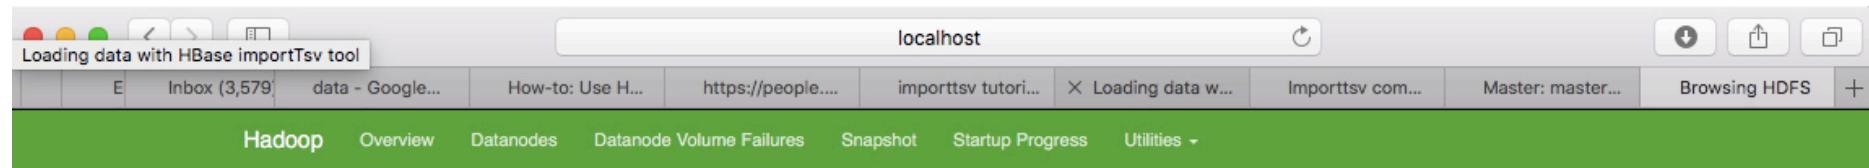
OSTech

## [Hbase – 2.4]

You can verify the file using the HDFS browser:

Access the HDFS web UI :<http://localhost:9870/>

Utilities -> Browse the file system -> /data



## Browse Directory

Browse Directory								
/data								
Show 25 entries								
□	Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
□	-rw-r--r--	root	supergroup	7.05 KB	May 04 16:13	2	128 MB	word_count.csv

First you need to design the table. To keep things simple, call it “wordcount” – the row keys will be the words themselves and the only column will contain the count, in a family that we’ll call “count”.

OSTech

## [Hbase – 2.4]

The best practice when creating a table is to split it according to the row key distribution but for this example we'll just create five regions with split points spread evenly across the key space. Open the hbase shell:

```
# hbase shell
```

And run the following command to create the table:

```
#create 'wordcount', {NAME => 'count'},  {SPLITS => ['g', 'm', 'r', 'w']}
```

```
[root@master0 data]# hbase shell
2022-05-04 10:48:23,098 WARN  [main] util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using built-in java classes where applicable
HBase Shell
Use "help" to get list of supported commands.
Use "exit" to quit this interactive shell.
For Reference, please visit: http://hbase.apache.org/2.0/book.html#shell
Version 2.4.11, r7e672a0da0586e6b7449310815182695bc6ae193, Tue Mar 15 10:31:00 PDT 2022
Took 0.0029 seconds
hbase:001:0> create 'wordcount', {NAME => 'count'},  {SPLITS => ['g', 'm', 'r', 'w']}
Created table wordcount
Took 7.5140 seconds
=> Hbase::Table - wordcount
hbase:002:0> █
```

## [Hbase – 2.4]

Verify the table:

<http://localhost:16010/tablesDetailed.jsp>

Click on Table Details → Wordcount

```
word  'wordcount', {NAME => 'count', BLOOMFILTER => 'ROW', IN_MEMORY => 'false', VERSIONS => '1', KEEP_DELETED_CELLS => 'FALSE', DATA_BLOCK_ENCODING => 'NONE', COMPRESSION => 'NONE', T
coun  TL => 'FOREVER', MIN_VERSIONS => '0', BLOCKCACHE => 'true', BLOCKSIZE => '65536', REPLICATION_SCOPE => '0'}
t
```

Determine the no of Regions: 5

### Table Regions

Name(5)	Region Server	ReadRequests (0)	WriteRequests (0)	StorefileSize (0 MB)	Num.Storefiles (0)	MemSize (0 MB)	Start Key	End Key	Region State
wordcount,,1651661388861.66b297fc6f59b59262b694088fa15913.	master0.hp.com:160 30	0	0	0 MB	0	0 MB	g		OPEN
wordcount,g,1651661388861.186db9219698cb4756d6a2371e9b7f37.	master0.hp.com:160 30	0	0	0 MB	0	0 MB	g	m	OPEN
wordcount,m,1651661388861.12c9c699e640d152ac628367e5e428b9.	master0.hp.com:160 30	0	0	0 MB	0	0 MB	m	r	OPEN
wordcount,r,1651661388861.37c4b40731617b8dd8005e4a805c6407.	master0.hp.com:160 30	0	0	0 MB	0	0 MB	r	w	OPEN
wordcount,w,1651661388861.d1bade65ab2f1ba4e788f087d0a1b0ab.	master0.hp.com:160 30	0	0	0 MB	0	0 MB	w		OPEN

## [Hbase – 2.4]

Invoking the HBase jar on the command line with the “hadoop” script will show a list of available tools.

```
# hbase org.apache.hadoop.hbase.mapreduce.ImportTsv
```

The one we want is called importtsv and has the following usage:

(format : hbase – Import Util – Columns - table – Importing file)

Open a New terminal and execute the following to generate StoreFiles for bulk-loading:

```
# hbase org.apache.hadoop.hbase.mapreduce.ImportTsv -Dimporttsv.separator=, -  
Dimporttsv.bulk.output=hdfs://mastero.hp.com:9000/data/hfile_temp -  
Dimporttsv.columns="HBASE_ROW_KEY,count:thevalue" wordcount  
hdfs://mastero.hp.com:9000/data/word_count.csv
```

## [Hbase – 2.4]

```
Reduce output records=722
Spilled Records=1444
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=115
CPU time spent (ms)=3920
Physical memory (bytes) snapshot=635961344
Virtual memory (bytes) snapshot=6834982912
Total committed heap usage (bytes)=442687488
ImportTsv
    Bad Lines=0
Shuffle Errors
    BAD_ID=0
    CONNECTION=0
    IO_ERROR=0
    WRONG_LENGTH=0
    WRONG_MAP=0
    WRONG_REDUCE=0
File Input Format Counters
    Bytes Read=7217
File Output Format Counters
    Bytes Written=36864
[root@master0 data]# hbase org.apache.hadoop.hbase.mapreduce.ImportTsv -Dimporttsv.separator=, -Dimporttsv.bulk.output=/data/hfile_temp -Dimporttsv.columns="HBASE_ROW_KEY, count:thevalue" wordcount /data/word_count.csv
```

Here's a rundown of the different configuration elements:

- **-Dimporttsv.separator=**, specifies that the separator is a comma.

## [Hbase – 2.4]

- **-Dimporttsv.bulk.output=output** is a relative path to where the HFiles will be written. the files will be in /data/. Skipping this option will make the job write directly to HBase.
- **-Dimporttsv.columns=HBASE\_ROW\_KEY,count** is a list of all the columns contained in this file. The row key needs to be identified using the all-caps HBASE\_ROW\_KEY string; otherwise it won't start the job. (I decided to use the qualifier "count" but it could be anything else.)

The job should complete within a minute, given the small input size. Note that five Reducers are running, one per region. Here's the result on HDFS:

Verify the files using the HDFS Browser:

## [Hbase – 2.4]

The screenshot shows the Apache HBase Master interface. At the top, there are several tabs: 'HBase Bul...', 'Apache HBase™ Refer...', 'Apache HBase™ Refer...', 'HBase Master: master0...', 'Edit Pad - Online Text E...', 'Browsing HDFS', and 'java - What is the'. Below these tabs is a green navigation bar with links: 'Hadoop', 'Overview', 'Datanodes', 'Datanode Volume Failures', 'Snapshot', 'Startup Progress', and 'Utilities'.

The main content area is titled 'Browse Directory' and displays the path '/data/hfile\_temp/count'. There are buttons for 'Go!', a folder icon, a file icon, and a refresh icon. Below this, there are filters for 'Show 25 entries' and a search bar.

A table lists five entries in the directory:

	Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
<input type="checkbox"/>	-rw-r--r--	root	supergroup	10.29 KB	May 04 20:17	3	128 MB	004b6bab9c242b7b06c77af36d33e36
<input type="checkbox"/>	-rw-r--r--	root	supergroup	15.95 KB	May 04 20:17	3	128 MB	6343b4e084334e03956b18d039483cd4
<input type="checkbox"/>	-rw-r--r--	root	supergroup	11.2 KB	May 04 20:17	3	128 MB	8dfbd3ce0d1f46069a1ee463d10ec8a2
<input type="checkbox"/>	-rw-r--r--	root	supergroup	6.49 KB	May 04 20:17	3	128 MB	a1e27782934a46dda0b0d5a6c7933605
<input type="checkbox"/>	-rw-r--r--	root	supergroup	12.7 KB	May 04 20:17	3	128 MB	c0b1d0e4d85c474495e3d6f8b80b6316

At the bottom, it says 'Showing 1 to 5 of 5 entries' and has buttons for 'Previous', '1', and 'Next'.

Using hdfs command.

```
#hdfs dfs -ls -R /data/hfile_temp
```

OSTech

## [Hbase – 2.4]

```
drwxr-xr-x - root supergroup          0 2022-05-04 14:47 /data/hfile_temp/count
[root@master0 hadoop]# hdfs dfs -ls -R /data/hfile_temp
2022-05-04 14:56:39,225 INFO Configuration.deprecation: No unit for dfs.client.datanode-restart.timeout(30) assuming SECONDS
-rw-r--r--  3 root supergroup          0 2022-05-04 14:47 /data/hfile_temp/_SUCCESS
drwxr-xr-x - root supergroup          0 2022-05-04 14:47 /data/hfile_temp/count
-rw-r--r--  3 root supergroup        10539 2022-05-04 14:47 /data/hfile_temp/count/004b6bab9c242b7b06c77af36d33e36
-rw-r--r--  3 root supergroup        16336 2022-05-04 14:47 /data/hfile_temp/count/6343b4e084334e03956b18d039483cd4
-rw-r--r--  3 root supergroup        11465 2022-05-04 14:47 /data/hfile_temp/count/8dfbd3ce0d1f46069a1ee463d10ec8a2
-rw-r--r--  3 root supergroup        6646  2022-05-04 14:47 /data/hfile_temp/count/a1e27782934a46dda0b0d5a6c7933605
-rw-r--r--  3 root supergroup        13009 2022-05-04 14:47 /data/hfile_temp/count/c0b1d0e4d85c474495e3d6f8b80b6316
[root@master0 hadoop]#
```

For the final step, we need to use the completebulkload tool to point to where the files are and which tables we are loading to:

```
#hbase org.apache.hadoop.hbase.mapreduce.LoadIncrementalHFiles
hdfs://master0.hp.com:9000/data/hfile_temp wordcount
```

- Argument : First – the Output folder specified in the importTSV output parameter.
- Second : Name of the Table.

## [Hbase – 2.4]

```
2022-05-04 15:00:20,824 INFO [ReadOnlyZKClient-127.0.0.1:2181@0x5be6e01c-SendThread(127.0.0.1:2181)] zookeeper.ClientCnxn: Opening socket connection to server localhost/127.0.0.1:2181. Will not attempt to authenticate using SASL (unknown error)
2022-05-04 15:00:20,859 INFO [ReadOnlyZKClient-127.0.0.1:2181@0x5be6e01c-SendThread(127.0.0.1:2181)] zookeeper.ClientCnxn: Socket connection established, initiating session, client: /127.0.0.1:59950, server: localhost/127.0.0.1:2181
2022-05-04 15:00:20,912 INFO [ReadOnlyZKClient-127.0.0.1:2181@0x5be6e01c-SendThread(127.0.0.1:2181)] zookeeper.ClientCnxn: Session establishment complete on server localhost/127.0.0.1:2181, sessionid = 0x100000e926e004c, negotiated timeout = 90000
2022-05-04 15:00:36,880 WARN [main] tool.LoadIncrementalHFiles: Skipping non-directory hdfs://master0.hp.com:9000/data/hfile_temp/_SUCCESS
2022-05-04 15:00:38,516 INFO [main] metrics.MetricRegistries: Loaded MetricRegistries class org.apache.hadoop.hbase.metrics.impl.MetricRegistriesImpl
2022-05-04 15:00:41,203 INFO [LoadIncrementalHFiles-0] tool.LoadIncrementalHFiles: Trying to load hfile=hdfs://master0.hp.com:9000/data/hfile_temp/count/004b6babcc9c242b7b06c77af36d33e36 first=Optional[m] last=Optional[quickstart]
2022-05-04 15:00:41,215 INFO [LoadIncrementalHFiles-1] tool.LoadIncrementalHFiles: Trying to load hfile=hdfs://master0.hp.com:9000/data/hfile_temp/count/6343b4e084334e03956b18d039483cd4 first=Optional[a] last=Optional[from]
2022-05-04 15:00:41,304 INFO [LoadIncrementalHFiles-1] tool.LoadIncrementalHFiles: Trying to load hfile=hdfs://master0.hp.com:9000/data/hfile_temp/count/8dfbd3ce0d1f46069a1ee463d10ec8a2 first=Optional[g] last=Optional[lot]
2022-05-04 15:00:41,367 INFO [LoadIncrementalHFiles-1] tool.LoadIncrementalHFiles: Trying to load hfile=hdfs://master0.hp.com:9000/data/hfile_temp/count/c0b1d0e4d85c474495e3d6f8b80b6316 first=Optional[r] last=Optional[vms]
2022-05-04 15:00:41,375 INFO [LoadIncrementalHFiles-0] tool.LoadIncrementalHFiles: Trying to load hfile=hdfs://master0.hp.com:9000/data/hfile_temp/count/a1e27782934a46dda0b0d5a6c7933605 first=Optional[w] last=Optional[yourself]
2022-05-04 15:00:42,504 INFO [main] client.ConnectionImplementation: Closing master protocol: MasterService
2022-05-04 15:00:42,628 INFO [ReadOnlyZKClient-127.0.0.1:2181@0x5be6e01c-EventThread] zookeeper.ClientCnxn: EventThread shut down for session : 0x100000e926e004c
2022-05-04 15:00:42,628 INFO [ReadOnlyZKClient-127.0.0.1:2181@0x5be6e01c] zookeeper.ZooKeeper: Session: 0x100000e926e004c closed
[root@master0 data]# hbase org.apache.hadoop.hbase.mapreduce.LoadIncrementalHFiles hdfs://master0.hp.com:9000/data/hfile_temp wordcount
```

Going back into the HBase shell, you can run the count command that will show you how many rows were loaded. If you forgot to chown, the command will hang.

```
# scan 'wordcount'
```

## [Hbase – 2.4]

```
words          column=count:thevalue, timestamp=2022-05-04T13:58:35.806, value=3
work           column=count:thevalue, timestamp=2022-05-04T13:58:35.806, value=1
workload        column=count:thevalue, timestamp=2022-05-04T13:58:35.806, value=2
works          column=count:thevalue, timestamp=2022-05-04T13:58:35.806, value=2
write           column=count:thevalue, timestamp=2022-05-04T13:58:35.806, value=11
writeheavy      column=count:thevalue, timestamp=2022-05-04T13:58:35.806, value=1
writing         column=count:thevalue, timestamp=2022-05-04T13:58:35.806, value=1
written         column=count:thevalue, timestamp=2022-05-04T13:58:35.806, value=3
wrong           column=count:thevalue, timestamp=2022-05-04T13:58:35.806, value=1
wrote           column=count:thevalue, timestamp=2022-05-04T13:58:35.806, value=1
you             column=count:thevalue, timestamp=2022-05-04T13:58:35.806, value=43
youll          column=count:thevalue, timestamp=2022-05-04T13:58:35.806, value=2
your            column=count:thevalue, timestamp=2022-05-04T13:58:35.806, value=24
youre          column=count:thevalue, timestamp=2022-05-04T13:58:35.806, value=2
yourself        column=count:thevalue, timestamp=2022-05-04T13:58:35.806, value=1
722 row(s)
Took 3.3041 seconds
hbase:008:0> scan 'wordcount'
```

You can verify the table regions details from the web UI:

<http://localhost:16010/master-status>

Home -> Table (User Tables) → Wordcount → Table regions:

## [Hbase – 2.4]

### Table Regions

Base Stats	Localities	Compactions	Name(5)	Region Server	ReadRequests (722)	WriteRequests (5)	StorefileSize (5 MB)	Num.Storefiles (5)	MemSize (0 MB)	Start Key	End Key	Region State
			wordcount,,1651672640769.0b49ac33cea190195653c14bbf52 bcb2.	master0.hp.com:160 30	246	1	1 MB	1	0 MB	g		OPEN
			wordcount,g,1651672640769.e502c2ac2a17e8f917f2ab4ebbf1 8533.	master0.hp.com:160 30	137	1	1 MB	1	0 MB	g	m	OPEN
			wordcount,m,1651672640769.07977d3dee2bc8bf9e11510211e c5dc4.	master0.hp.com:160 30	121	1	1 MB	1	0 MB	m	r	OPEN
			wordcount,r,1651672640769.9bb89b702a505d94d460baf6264f ec5e.	master0.hp.com:160 30	180	1	1 MB	1	0 MB	r	w	OPEN
			wordcount,w,1651672640769.ce6fad66a0aa9ac9718233fb93c7 785c.	master0.hp.com:160 30	38	1	1 MB	1	0 MB	w		OPEN

As you can view from above, all the Memsize values are Zero. This means the records are not gone through the normal process.

Insert a record using CLI and verify the record:

```
# put 'wordcount' , 'henry' , 'count:thevalue' , 15
# scan 'wordcount'
```

After that verify the Table Regions:

OSTech

## [Hbase – 2.4]

### Table Regions

Base Stats	Localities	Compactions										
Name(5)	Region Server	ReadRequests (2,167)	WriteRequests (6)	StorefileSize (5 MB)	Num.Storefiles (5)	MemSize (1 MB)	Start Key	End Key	Region State			
wordcount,,1651672640769.0b49ac33cea190195653c14bbf52 bcb2.	master0.hp.com:16030	738	1	1 MB	1	0 MB	g		OPEN			
wordcount,g,1651672640769.e502c2ac2a17e8f917f2ab4ebbf1 8533.	master0.hp.com:16030	412	2	1 MB	1	1 MB	g	m	OPEN			
wordcount,m,1651672640769.07977d3dee2bc8bf9e11510211e c5dc4.	master0.hp.com:16030	363	1	1 MB	1	0 MB	m	r	OPEN			
wordcount,r,1651672640769.9bb89b702a505d94d460baf6264f ec5e.	master0.hp.com:16030	540	1	1 MB	1	0 MB	r	w	OPEN			
wordcount,w,1651672640769.ce6fad66a0aa9ac9718233fb93c7 785c.	master0.hp.com:16030	114	1	1 MB	1	0 MB	w		OPEN			

As you can see, the memtable get populated.

----- Lab Ends Here -----

OSTech

## 11. Hbase Spark Integration – Spark-Hbase connector – 90 Minutes

Compatible: Hbase 2.4 and Spark – 3.0.2

Stop the Hbase cluster and revert to hbase standalone mode. Refer the Annexure (Standalone Mode)

We will perform the following in this lab in a standalone Hbase installation:

- Configure Spark to intergrate with Hbase.
- Retrieve records from Hbase table using Spark and create dataframe.
- Store spark dataframe in Hbase table.

Extract hbase connector and then copy the hbase\* jar from the lib folder to the spark home jars folder.

```
#cd /Software/hbase-connectors-1.0.1-SNAPSHOT/lib  
#cp hbase-* /opt/spark/jars/
```

```
[root@master0 lib]# ls hbase*  
hbase-annotations-2.2.4-tests.jar    hbase-kafka-proxy-1.0.1-SNAPSHOT.jar  hbase-shaded-netty-2.2.1.jar  
hbase-client-2.2.4.jar               hbase-metrics-2.2.4.jar          hbase-shaded-protobuf-2.1.0.jar  
hbase-common-2.2.4.jar              hbase-metrics-api-2.2.4.jar     hbase-shaded-protobuf-2.2.1.jar  
hbase-hadoop-compat-2.2.4.jar      hbase-protocol-2.2.4.jar        hbase-spark-1.0.1-SNAPSHOT.jar  
hbase-hadoop2-compat-2.2.4.jar     hbase-protocol-shaded-2.2.4.jar   hbase-spark-protocol-1.0.1-SNAPSHOT.jar  
hbase-kafka-model-1.0.1-SNAPSHOT.jar hbase-shaded-miscellaneous-2.2.1.jar hbase-spark-protocol-shaded-1.0.1-SNAPSHOT.jar  
[root@master0 lib]# pwd  
/Software/hbase-connectors-1.0.1-SNAPSHOT/lib  
[root@master0 lib]#
```

## [Hbase – 2.4]

Extract Phoenix to a folder and copy the phoenix related jar to the spark jar folder. (Assuming that the phoenix is extracted in /opt/phoenix folder)

Copy Phoenix jar, it contains all the dependencies jar required for interaction with spark and hbase.

```
#cd /opt/phoenix  
#cp phoenix-* /opt/spark/jars/
```

```
[root@master0 jars]# ls ph*  
phoenix-client-hbase-2.4.0-5.1.2.jar  phoenix-phref-5.1.2.jar  phoenix-server-hbase-2.4.0-5.1.2.jar  
[root@master0 jars]#
```

Prepare HBase table with data

Run the following commands in HBase shell to prepare a sample table that will be used in the following sections.

```
# hbase shell
```

```
create 'Person', 'Name', 'Address'  
put 'Person', '1', 'Name:First', 'Raymond'  
put 'Person', '1', 'Name:Last', 'Tang'  
put 'Person', '1', 'Address:Country', 'Australia'  
put 'Person', '1', 'Address:State', 'VIC'
```

```
put 'Person', '2', 'Name:First', 'Dnomyar'  
put 'Person', '2', 'Name:Last', 'Gnat'  
put 'Person', '2', 'Address:Country', 'USA'
```

## [Hbase – 2.4]

```
put 'Person', '2', 'Address:State', 'CA'
```

The table returns the following result when scanning:

```
scan 'Person'
```

**Case 1 :** Fetch records from a Hbase and Store in a file using Spark.

Open a terminal and start a spark shell.

```
#spark-shell
```

Create DataFrame with the following commands. Execute all the commands in spark shell.

1) First import the required classes:

```
import org.apache.hadoop.hbase.spark.HBaseContext  
import org.apache.hadoop.hbase.HBaseConfiguration
```

2) Create HBase configurations

```
val conf = HBaseConfiguration.create()
```

OSTech

## [Hbase – 2.4]

```
conf.set("hbase.zookeeper.quorum", "127.0.0.1:2181")
```

### 3) Create HBase context

```
// Instantiate HBaseContext that will be used by the following code  
new HBaseContext(spark.sparkContext, conf)
```

### 4) Create DataFrame

```
val hbaseDF = (spark.read.format("org.apache.hadoop.hbase.spark")  
.option("hbase.columns.mapping",  
"rowKey STRING :key," +  
"firstName STRING Name:First, lastName STRING Name>Last," +  
"country STRING Address:Country, state STRING Address:State"  
)  
.option("hbase.table", "Person")  
.load()
```

The columns mapping matches with the definition in the steps above.

### 5) Show DataFrame schema

```
scala> hbaseDF.schema
```

### 6) Show data

```
hbaseDF.show()
```

OSTech

## [Hbase – 2.4]

```
scala> val hbaseDF = (spark.read.format("org.apache.hadoop.hbase.spark")
|   .option("hbase.columns.mapping",
|     "rowKey STRING :key," +
|     "firstName STRING Name:First, lastName STRING Name:Last," +
|     "country STRING Address:Country, state STRING Address:State"
|   )
|   .option("hbase.table", "Person")
|   ).load()
rowKey STRING :key,firstName STRING Name:First, lastName STRING Name:Last,country STRING Address:Country, state STRING
Address:State
hbaseDF: org.apache.spark.sql.DataFrame = [lastName: string, country: string ... 3 more fields]

scala> hbaseDF.schema
res2: org.apache.spark.sql.types.StructType = StructType(StructField(lastName,StringType,true), StructField(country,St
ringType,true), StructField(state,StringType,true), StructField(firstName,StringType,true), StructField(rowKey,StringT
ype,true))

scala> hbaseDF.show()
+-----+-----+-----+-----+
|lastName|country|state|firstName|rowKey|
+-----+-----+-----+-----+
|    Tang|Australia|  VIC|  Raymond|     1|
|    Gnat|        USA|   CA| Dnomyar|     2|
+-----+-----+-----+-----+
```

Until now, we've successfully loaded data from HBase in Spark 3.0.1.

You can also write into HBase from Spark too

OSTech

[Hbase – 2.4]

## Case 2 : Fetching records from a file/from Memory dataframe and insert into Hbase using Spark.

In this example we want to store personal data in an HBase table. We want to store name as string, email address as string, birth date as date and height as a floating point number. The contact information (email) is stored in the **c** column family and personal information (birth date, height) is stored in the **p** column family. The key in HBase table will be the **name** attribute.

	Spark	HBase
Type/Table	Person	person
Name	name: String	key
Email address	email: String	c:email
Birth date	birthDate: Date	p:birthDate
Height	height: Float	p:height

Create HBase table – Using hbase shell.

Use the following command to create the HBase table:

```
shell> create 'person', 'p', 'c'
```

## [Hbase – 2.4]

Insert data – In Spark-shell.

Use the following spark code in spark-shell to insert data into our HBase table:

```
val sql = spark.sqlContext  
  
import java.sql.Date  
  
case class Person(name: String,  
                  email: String,  
                  birthDate: Date,  
                  height: Float)  
  
var personDS = Seq(  
  Person("alice", "alice@alice.com", Date.valueOf("2000-01-01"), 4.5f),  
  Person("bob", "bob@bob.com", Date.valueOf("2001-10-17"), 5.1f)  
) .toDS
```

[Hbase – 2.4]

```
scala> val sql = spark.sqlContext
sql: org.apache.spark.sql.SQLContext = org.apache.spark.sql.SQLContext@1cc784c9

scala> import java.sql.Date
import java.sql.Date

scala>

scala> case class Person(name: String,
|                     email: String,
|                     birthDate: Date,
|                     height: Float)
defined class Person

scala>

scala> var personDS = Seq(
|   Person("alice", "alice@alice.com", Date.valueOf("2000-01-01"), 4.5f),
|   Person("bob", "bob@bob.com", Date.valueOf("2001-10-17"), 5.1f)
| ).toDS
personDS: org.apache.spark.sql.Dataset[Person] = [name: string, email: string ... 2 more fields]

scala> █
```

OSTech

## [Hbase – 2.4]

Save the datframe to a Hbase table.

```
personDS.write.format("org.apache.hadoop.hbase.spark")
.option("hbase.columns.mapping",
  "name STRING :key, email STRING c:email, " +
  "birthDate DATE p:birthDate, height FLOAT p:height")
.option("hbase.table", "person")
.option("hbase.spark.use.hbasecontext", false)
.save()
```

The previously inserted data can be tested with a simple scan: On the Hbase shell.

```
shell> scan 'person'
```

## [Hbase – 2.4]

```
Use "exit" to quit this interactive shell.
For Reference, please visit: http://hbase.apache.org/2.0/book.html#shell
Version 2.4.11, r7e672a0da0586e6b7449310815182695bc6ae193, Tue Mar 15 10:31:00 PDT 2022
Took 0.0065 seconds
hbase:001:0> scan 'person'
ROW                                COLUMN+CELL
alice                             column=c:email, timestamp=2022-05-06T15:58:19.976, value=alice@alice.com
alice                             column=p:birthDate, timestamp=2022-05-06T15:58:19.976, value=\x00\x00\x00\xDCj\xCF\xAC\
                           x00
alice                             column=p:height, timestamp=2022-05-06T15:58:19.976, value=@\x90\x00\x00
bob                               column=c:email, timestamp=2022-05-06T15:58:19.990, value=bob@bob.com
bob                               column=p:birthDate, timestamp=2022-05-06T15:58:19.990, value=\x00\x00\x00\xE9\x97\xF5\x
                           10\x00
bob                               column=p:height, timestamp=2022-05-06T15:58:19.990, value=@\xA333
2 row(s)
```

You have successfully inserted record to hbase table using Spark Dataframe.

-----Lab Ends Here -----

## 12. Install MySQL on a CentOS 7 Server

```
wget https://dev.mysql.com/get/mysql80-community-release-el7-3.noarch.rpm  
rpm -Uvh mysql80-community-release-el7-3.noarch.rpm  
yum install mysql-server
```

### Start MySQL and Check its Status

- Once you have MySQL ready on CentOS 7, it does not automatically start right after the installation. Therefore, you need to start it manually through the following command:

```
systemctl start mysqld
```

- You will get no response once MySQL starts so to check if it is working properly, use the command below:

```
systemctl status mysqld
```

When installing MySQL on CentOS 7, a temporary root password is generated. Issue the command below to see it:  
`grep 'password' /var/log/mysqld.log`

In order to change it, follow these steps:

- Firstly, run the following command:

```
sudo mysql_secure_installation  
password: Root123.
```

CLI

OSTech

## [Hbase – 2.4]

```
#mysql -u root -p  
##### Lab Ends Here -----  
Source: https://www.hostinger.in/tutorials/how-to-install-mysql-on-centos-7
```

### 13. Errors:

**ERROR:** org.apache.hadoop.hbase.PeerHoldException: Master is initializing  
Scenario: Unable to create table using the hbase shell.

**Solution :** clean /opt/hbase/tmp and restart hbase server else try the following commands.

Install zookeeper use its library to access hbase's zookeeper to remove the following entry.

```
# tar -xvf apache-zookeeper* -c /opt  
# cd /opt/apache-zookeeper*  
# bin/zkCli.sh  
[zk: localhost:2181(CONNECTED) 3] get /hbase/meta-region-server  
?master:16000?m l5???PBUF
```

mastero.hp.com?}?????o

```
#delete /hbase/meta-region-server
```

OSTech

[Hbase – 2.4]

Restart the hbase.

Error: Unable to access from Laptop in docker.

```
127.0.0.1 Henrys-MacBook-Air.local localhost
255.255.255.255 broadcasthost
::1           localhost
127.0.0.1 couchbase0 master0 master0.hp.com
# Added by Docker Desktop
# To allow the same kube context to work on the host and the container:
127.0.0.1 kubernetes.docker.internal
# End of section
```

Ensure to enter the hostname of the Hbase Node in the Laptop /etc/hosts

Error

```
ERROR: org.apache.hadoop.hbase.PeerInitializationException: Master is initializing
      at org.apache.hadoop.hbase.master.HMaster.checkInitialized(HMaster.java:2829)
      at org.apache.hadoop.hbase.master.HMaster.createTable(HMaster.java:2085)
      at
org.apache.hadoop.hbase.master.MasterRpcServices.createTable(MasterRpcServices.java:706)
```

OSTech

## [Hbase – 2.4]

at

```
org.apache.hadoop.hbase.shaded.protobuf.generated.MasterProtos$MasterService$2.callBlocking
Method(MasterProtos.java)
    at org.apache.hadoop.hbase.ipc.RpcServer.call(RpcServer.java:392)
    at org.apache.hadoop.hbase.ipc.CallRunner.run(CallRunner.java:133)
    at org.apache.hadoop.hbase.ipc.RpcExecutor$Handler.run(RpcExecutor.java:359)
    at org.apache.hadoop.hbase.ipc.RpcExecutor$Handler.run(RpcExecutor.java:339)
```

Solution: Clean the zookeeper directory.

<https://stackoverflow.com/questions/17038957/org-apache-hadoop-hbase-pleaseholdexception-master-is-initializing>

**Error: java.io.IOException: cannot get log writer**

```
2022-06-04 02:32:15,842 INFO [master/mastero:16000:becomeActiveMaster] wal.AbstractFSWAL: Closed WAL: AsyncFSWAL
mastero.hp.com%2C16000%2C1654309909250:(num 1654309935646)
2022-06-04 02:32:15,917 ERROR [master/mastero:16000:becomeActiveMaster] master.HMaster: Failed to become active master
java.io.IOException: cannot get log writer
```

```
    at org.apache.hadoop.hbase.wal.AsyncFSWALProvider.createAsyncWriter(AsyncFSWALProvider.java:128)
    at org.apache.hadoop.hbase.regionserver.wal.AsyncFSWAL.createWriterInstance(AsyncFSWAL.java:689)
    at org.apache.hadoop.hbase.regionserver.wal.AsyncFSWAL.createWriterInstance(AsyncFSWAL.java:130)
    at org.apache.hadoop.hbase.regionserver.wal.AbstractFSWAL.rollWriter(AbstractFSWAL.java:841)
    at org.apache.hadoop.hbase.regionserver.wal.AbstractFSWAL.rollWriter(AbstractFSWAL.java:548)
    at org.apache.hadoop.hbase.regionserver.wal.AbstractFSWAL.init(AbstractFSWAL.java:489)
    at org.apache.hadoop.hbase.wal.AsyncFSWALProvider.getWAL(AsyncFSWALProvider.java:160)
    at org.apache.hadoop.hbase.wal.AsyncFSWALProvider.getWAL(AsyncFSWALProvider.java:62)
    at org.apache.hadoop.hbase.wal.WALFactory.getWAL(WALFactory.java:296)
```

OSTech

## [Hbase – 2.4]

```
at org.apache.hadoop.hbase.master.region.MasterRegion.createWAL(MasterRegion.java:187)
at org.apache.hadoop.hbase.master.region.MasterRegion.bootstrap(MasterRegion.java:207)
```

Solution : hbase.unsafe.stream.capability.enforce to false in base-site.xml

**ERROR:** KeeperErrorCode = NoNode for /hbase/master

Solution: Check the log. Determine the Namenode port no etc.

#### **14. Annexure:**

Hbase & YARN – with Mounted Folder and Extra port.

```
#docker run --name hbaseo --hostname hmaster.master.com --privileged --network spark-net -v  
/Users/henrypotsangbam/Documents/Docker:/opt -p 2181:2181 -p 16000:16000 -p 16010:16010 -p  
16002:16002 -p 16012:16012 -p 16020:16020 -p 16030:16030 -p 8080:8080 -p 8085:8085 -p 8088:8088 -p  
9870:9870 -p 9864:9864 -p 9000:9000 -p 9001:9002 -p 9003:9003 -p 9004:9004 -p 9005:9005 -i -t centos:7  
/usr/sbin/init
```

```
#docker run --name monitor --hostname monitor --privileged --network spark-net -v  
/Users/henrypotsangbam/Documents/Docker:/opt -p 9090:9090 -p 9010:9010 -p 3000:3000 -i -t  
centos:7 /usr/sbin/init
```

## 15. Annexure (Standalone Mode)

Take a back up of hbase-site.xml. Location : /opt/hbase/conf

Replace the above file with the following content.

```
<?xml version="1.0"?>
<?xml-stylesheet type="text/xsl" href="configuration.xsl"?>
<configuration>
  <property>
    <name>hbase.rootdir</name>
    <value>file:///opt/data/hbase</value>
  </property>
  <property>
    <name>hbase.zookeeper.property.dataDir</name>
    <value>/opt/data/zookeeper</value>
  </property>
  <property>
    <name>hbase.unsafe.stream.capability.enforce</name>
    <value>false</value>
  </property>
</configuration>
```

Start the hbase standalone cluster: bin/start-hbase.sh